



# 第8章 异常控制流

## 异常和进程 Exceptions and Processes

100076202: 计算机系统导论

任课教师:

宿红毅 张艳 黎有琦 颜珂

原作者:


Randal E. Bryant and David R. O'Hallaron



**Carnegie  
Mellon  
University**

# 打印机过去常常着火

## Printers Used to Catch on Fire



# 高度异常控制流

## Highly Exceptional Control Flow



```
234 static int lp_check_status(int minor)
235 {
236     int error = 0;
237     unsigned int last = lp_table[minor].last_error;
238     unsigned char status = r_str(minor);
239     if ((status & LP_PERRORP) && !(LP_F(minor) & LP_CAREFUL))
240         /* No error. */
241         last = 0;
242     else if ((status & LP_POUTPA)) {
243         if (last != LP_POUTPA) {
244             last = LP_POUTPA;
245             printk(KERN_INFO "lp%d out of paper\n", minor);
246         }
247         error = -ENOSPC;
248     } else if (!(status & LP_PSELECD)) {
249         if (last != LP_PSELECD) {
250             last = LP_PSELECD;
251             printk(KERN_INFO "lp%d off-line\n", minor);
252         }
253         error = -EIO;
254     } else if (!(status & LP_PERRORP)) {
255         if (last != LP_PERRORP) {
256             last = LP_PERRORP;
257             printk(KERN_INFO "lp%d on fire\n", minor);
258         }
259         error = -EIO;
260     } else {
261         last = 0; /* Come here if LP_CAREFUL is set and no
262                  errors are reported. */
263     }
264
265     lp_table[minor].last_error = last;
266
267     if (last != 0)
268         lp_error(minor);
269
270     return error;
271 }
```

<https://git.kernel.org/pub/scm/linux/kernel/git/torvalds/linux.git/tree/drivers/char/lp.c?h=v5.0-rc3>



# 内容提纲

- **异常控制流** Exceptional Control Flow CSAPP 8
- **异常** Exceptions CSAPP 8.1
- **进程** Processes CSAPP 8.2
- **进程控制** Process Control CSAPP 8.3-8.4

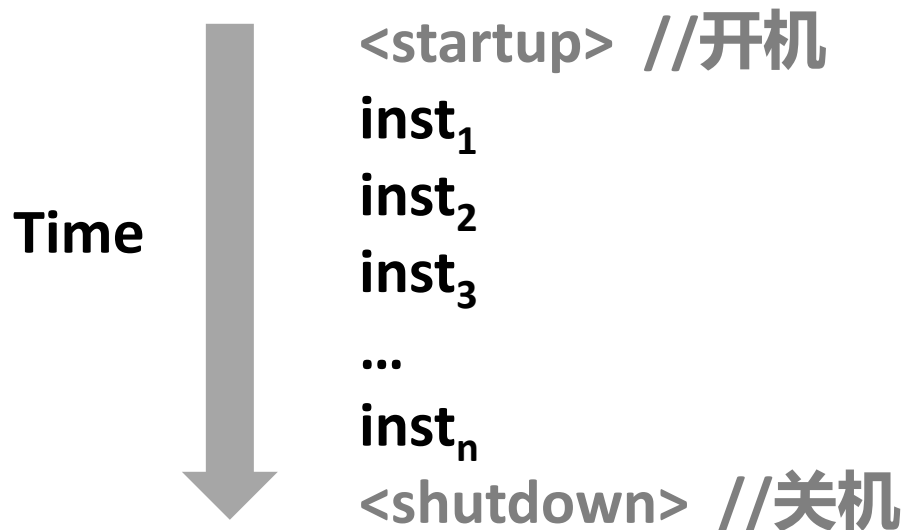
# 控制流 Control Flow



## ■ 处理器只做一件事 Processors do only one thing:

- 从开机到关机，每个CPU核只是读入和执行（解释）指令序列，每次一条 From startup to shutdown, each CPU core simply reads and executes (interprets) a sequence of instructions, one at a time \*
- 这个序列就是CPU的**控制流** This sequence is the CPU's *control flow* (or *flow of control*)

### **物理控制流** Physical control flow



\* 从外部体系结构视角看（内部来看，CPU可以使用并行乱序执行）

\* Externally, from an architectural viewpoint (internally, the CPU may use parallel out-of-order execution)

# 改变控制流 Altering the Control Flow



- **目前：两种改变控制流的机制：** Up to now: two mechanisms for changing control flow:
  - 跳转分支指令 Jumps and branches
  - 调用和返回指令 Call and return反应**程序状态**的变化 React to changes in **program state**
- **对有用的系统来说还不够：** Insufficient for a useful system:  
**难以反应系统状态的改变** Difficult to react to changes in **system state**
  - 从磁盘或者网络适配器获取的数据到达 Data arrives from a disk or a network adapter
  - 指令除零 Instruction divides by zero
  - 用户键盘按下了Ctrl-C **User hits Ctrl-C at the keyboard**
  - 系统定时器超时 System timer expires
- **系统需要“异常控制流”处理机制** System needs mechanisms for “exceptional control flow”

# 异常控制流 Exceptional Control Flow



- 存在计算机系统的每个层次 Exists at all levels of a computer system
- 低层次机制 Low level mechanisms
  - 1. 异常 Exceptions
    - 为响应系统事件改变控制流（例如系统状态改变） Change in control flow in response to a system event (i.e., change in system state)
    - 硬件和OS软件组合实现 Implemented using combination of hardware and OS software
- 高层次机制 Higher level mechanisms
  - 2. 进程上下文切换 Process context switch
    - 硬件定时器和OS软件实现 Implemented by OS software and hardware timer
  - 3. 信号 Signals
    - OS软件实现 Implemented by OS software
  - 4. 非局部跳转 Nonlocal jumps: `setjmp()` and `longjmp()`
    - C运行时库实现 Implemented by C runtime library





# 内容提纲

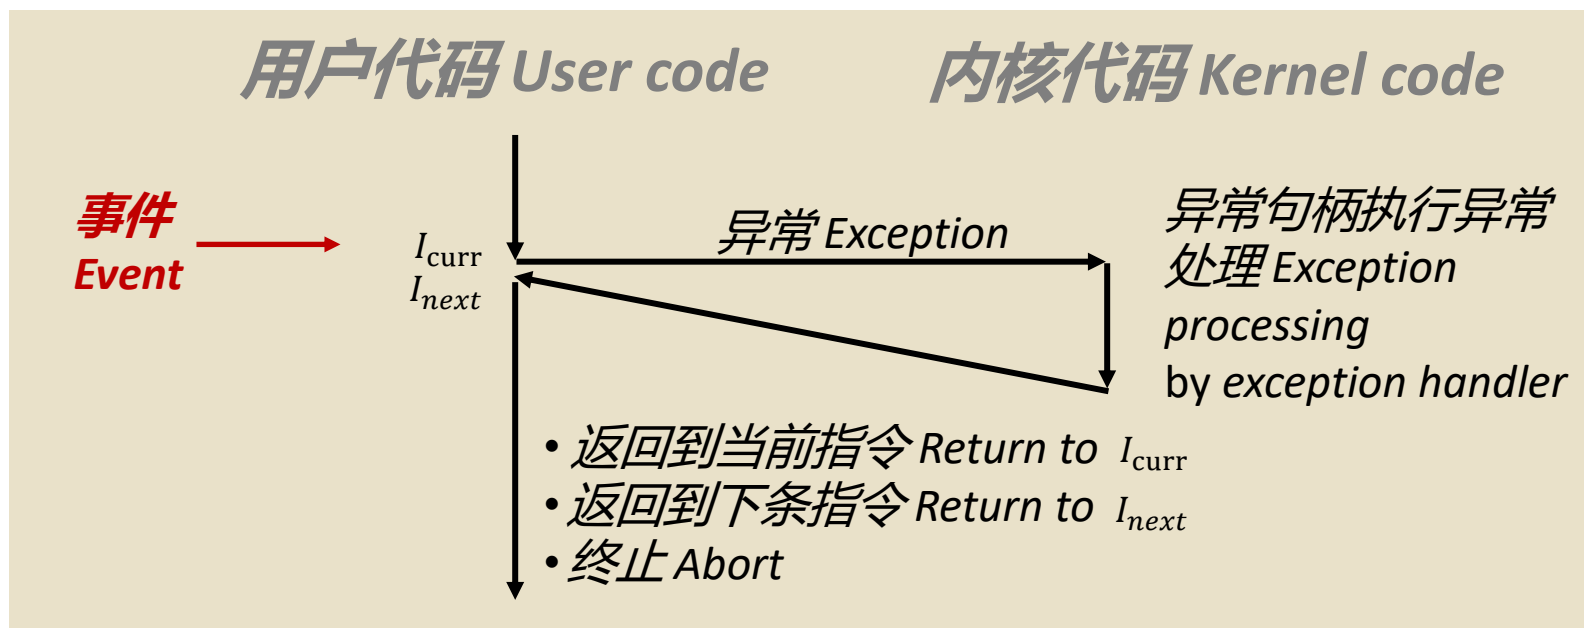
- 异常控制流 Exceptional Control Flow
- 异常 Exceptions
- 进程 Processes
- 进程控制 Process Control



# 异常 Exceptions



- **异常**是为了响应某些事件（即处理器状态改变）而将控制流转移到OS内核 An **exception** is a transfer of control to the OS *kernel* in response to some *event* (i.e., change in processor state)
  - 内核是操作系统的内存驻留 Kernel is the memory-resident part of the OS
  - 事件举例：除零，算术溢出，缺页，I/O请求完成，键入Ctrl+C Examples of events: Divide by 0, arithmetic overflow, page fault, I/O request completes, typing Ctrl-C

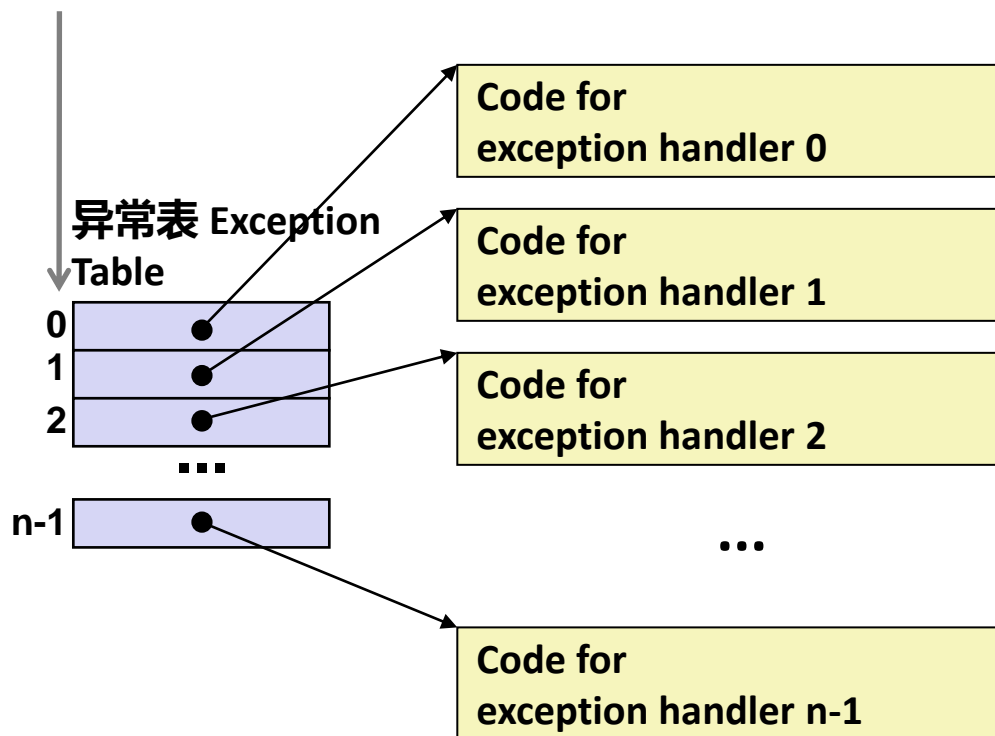




# 异常表 Exception Tables

异常号

Exception  
numbers



- 每个事件类型有惟一的异常编号k Each type of event has a unique exception number k
- 用k做为异常表的索引（即中断向量） k = index into exception table (a.k.a. interrupt vector)
- 每次发生异常k时，就会调用句柄k（句柄就是异常处理程序指针） Handler k is called each time exception k occurs

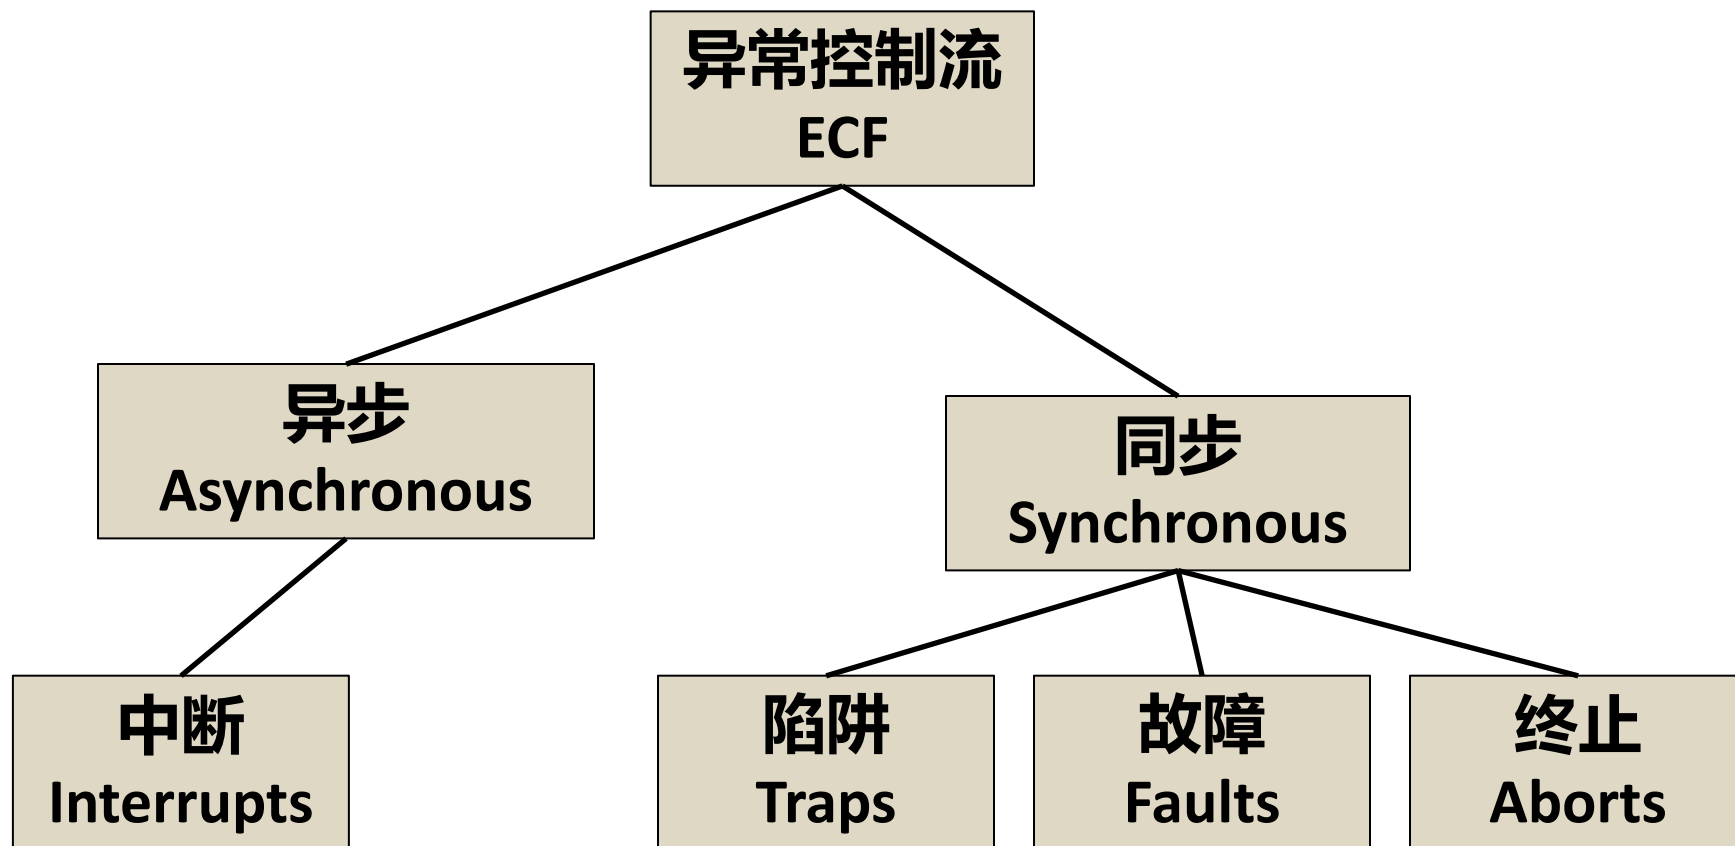


图 8-3 生成异常处理程序的地址。异常号是到异常表中的索引



# (部分) 异常分类 (partial)

## Taxonomy





# 异步异常（中断）

## Asynchronous Exceptions (Interrupts)

- **由处理器外部事件引起** Caused by events external to the processor
  - 通过设置处理器的中断引脚来指示有中断请求到达 Indicated by setting the processor's *interrupt pin*
  - 中断处理程序返回后执行下一条指令 Handler returns to “next” instruction
- **举例 Examples:**
  - 时钟中断 Timer interrupt
    - 每隔大约几ms，外部时钟芯片触发一个中断 Every few ms, an external timer chip triggers an interrupt
    - 将控制权从用户程序切换到内核 Used by the kernel to take back control from user programs
  - 外部设备的I/O中断 I/O interrupt from external device
    - 键盘键入Ctrl-C **Hitting Ctrl-C at the keyboard**
    - 网络有一个包抵达 Arrival of a packet from a network
    - 从磁盘有数据抵达 Arrival of data from a disk

# 同步异常 Synchronous Exceptions



## ■ 指令执行结果导致的异常事件 Caused by events that occur as a result of executing an instruction:

### ■ 陷入/陷阱 *Traps*

- 人为的 Intentional
- 例如：系统调用、断点、特殊指令等 Examples: *system calls*, breakpoint traps, special instructions
- 控制流返回到下一条指令 Returns control to “next” instruction

### ■ 故障 *Faults*

- 不是有意的但是大概率可恢复 Unintentional but possibly recoverable
- 例如：缺页异常（可恢复）、保护异常（不可恢复）、浮点异常 Examples: page faults (recoverable), protection faults (unrecoverable), floating point exceptions
- 重新执行故障（“当前”）指令或者终止执行 Either re-executes faulting (“current”) instruction or aborts

### ■ 终止 *Aborts*

- 非故意且不可恢复 Unintentional and unrecoverable
- 例如：非法指令、校验错、机器检查 Examples: illegal instruction, parity error, machine check
- 终止当前程序执行 Aborts current program



# 系统功能调用 System Calls

- 每个x86-64系统调用都有一个唯一的ID编号 Each x86-64 system call has a unique ID number
- 例如: Examples:

编号 <i>Number</i>	名字 <i>Name</i>	描述 <i>Description</i>
0	read	读文件 Read file
1	write	写文件 Write file
2	open	打开文件 Open file
3	close	关闭文件 Close file
4	stat	获取有关文件的信息 Get info about file
57	fork	创建进程 Create process
59	execve	执行一个程序 Execute a program
60	_exit	终止进程 Terminate process
62	kill	发送信号给进程 Send signal to process

# 系统调用举例：打开文件



## System Call Example: Opening File

- 用户调用：open函数 User calls: `open(filename, options)`
- 调用\_\_open函数，该函数会触发syscall系统功能调用指令 Calls `__open` function, which invokes system call instruction `syscall`

```
0000000000e5d70 <__open>:
```

```
...
```


```
e5d79: b8 02 00 00 00  mov $0x2,%eax # open is syscall #2
```

```
e5d7e: 0f 05          syscall      # Return value in %rax
```

```
e5d80: 48 3d 01 f0 ff ff  cmp $0xfffffffffff001,%rax
```

```
...
```

```
e5dfa: c3            retq
```



- `%rax`包含系统调用号 `%rax` contains syscall number
- 其它参数存放在 Other arguments in `%rdi`, `%rsi`, `%rdx`, `%r10`, `%r8`, `%r9`
- 返回值在 Return value in `%rax`
- 负值是错误号 Negative value is an error corresponding to negative `errno`



# 系统调用举例：打开文件



## System Call Example

- 用户调用：open
- 调用\_\_open函数，function, which in

```
00000000000e5d70 <__open@libc.so.6>
...
e5d79: b8 02 00 00 00
e5d7e: 0f 05          sys
e5d80: 48 3d 01 f0 ff ff
...
e5dfa: c3            retq
```

用户代码  
User code

syscall  
cmp

异常

返回

几乎和函数调用类似 Almost like a function call

- 转换控制 Transfer of control
- 返回时执行下条指令 On return, executes next instruction
- 使用调用规则传递参数 Passes arguments using calling convention
- 返回值在%rax中 Gets result in %rax

一个重要的差异 One Important exception!

- 由内核执行 Executed by Kernel
- 不同的优先权 Different set of privileges
- 以及其它不同：And other differences:
  - 例如：“函数”的“地址”是在%rax中 E.g., “address” of “function” is in %rax
  - 使用错误号 Uses errno
  - 等 Etc.

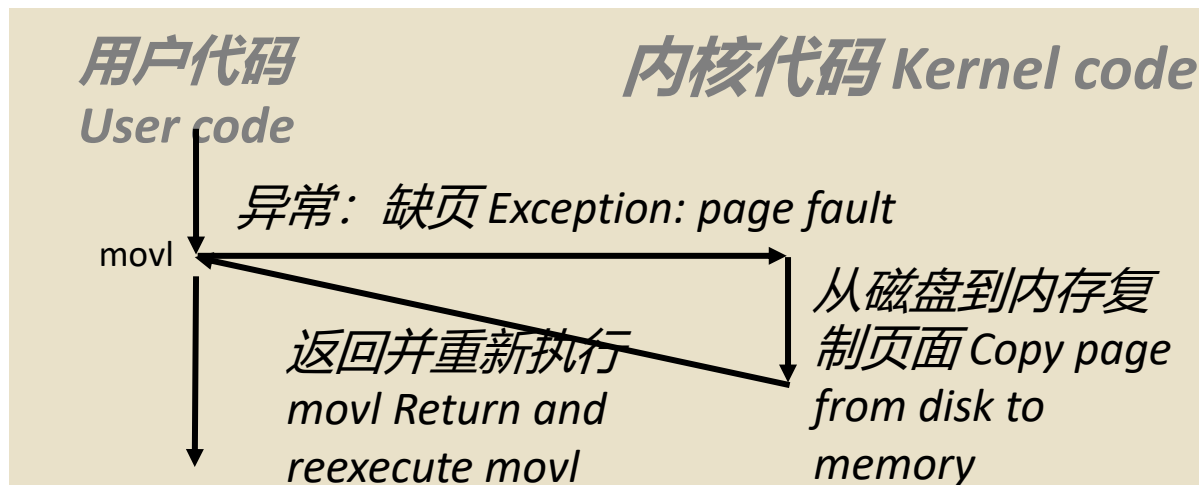


# 故障举例：缺页异常 Fault Example: Page Fault

- 用户写内存 User writes to memory location
- 对应的页面在磁盘上 That portion (page) of user's memory is currently on disk

```
int a[1000];  
main ()  
{  
    a[500] = 13;  
}
```

80483b7: c7 05 10 9d 04 08 0d movl \$0xd,0x8049d10





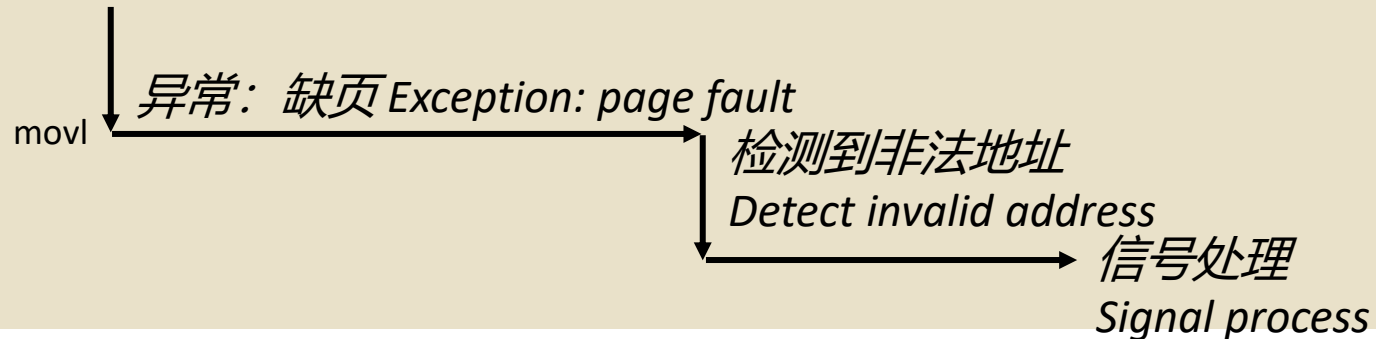
# 故障举例：非法内存引用

## Fault Example: Invalid Memory Reference

```
int a[1000];  
main ()  
{  
    a[5000] = 13;  
}
```

80483b7: c7 05 60 e3 04 08 0d movl \$0xd,0x804e360

用户代码 User code    内核代码 Kernel code



- 发送SIGSEGV信号给用户进程 Sends **SIGSEGV** signal to user process
- 用户进程会“段错误”异常退出 User process exits with “segmentation fault”



# 内容提纲

- 异常控制流 Exceptional Control Flow
- 异常 Exceptions
- 进程 Processes
- 进程控制 Process Control



# 进程 Processes


- 定义: **进程**是程序的一次执行(运行程序的实例)

Definition: A **process** is an instance of a running program.


- 计算机科学最重要的概念之一 One of the most profound ideas in computer science
- 与“程序”或“处理器”不同 Not the same as “program” or “processor”

- 进程为每个程序提供了两个关键抽象 Process provides each program with two key abstractions:

- **逻辑控制流** *Logical control flow*
  - 每个程序看起来独占CPU Each program seems to have exclusive use of the CPU
  - 内核支持的上下文切换 Provided by kernel mechanism called *context switching*
- **私有地址空间** *Private address space*
  - 每个程序看起来独占主存空间 Each program seems to have exclusive use of main memory.
  - 内核支持的虚拟内存 Provided by kernel mechanism called *virtual memory*



# 多进程幻象： Multiprocessing: The Illusion



- **计算机同时运行很多进程** Computer runs many processes simultaneously
  - 单个或多个用户的应用 Applications for one or more users
    - Web浏览器、邮件客户、编辑器。。。 Web browsers, email clients, editors, ...
  - 后台任务 Background tasks
    - 监视网络和I/O设备 Monitoring network & I/O devices

# 多进程举例 Multiprocessing Example



```
Processes: 123 total, 5 running, 9 stuck, 109 sleeping, 611 threads
Load Avg: 1.03, 1.13, 1.14 CPU usage: 3.27% user, 5.15% sys, 91.56% idle
SharedLibs: 576K resident, 0B data, 0B linkedit.
MemRegions: 27958 total, 1127M resident, 35M private, 494M shared.
PhysMem: 1039M wired, 1974M active, 1062M inactive, 4076M used, 18M free.
VM: 280G vsize, 1091M framework vsize, 23075213(1) pageins, 5843367(0) pageouts.
Networks: packets: 41046228/11G in, 66083096/77G out.
Disks: 17874391/349G read, 12847373/594G written.


PID    COMMAND    %CPU TIME    #TH    #WQ    #PORT    #MREG    RPRVT    RSHRD    RSIZE    VPRVT    VSIZE
99217-  Microsoft Of 0.0 02:28.34 4      1      202     418     21M     24M     21M     66M     763M
99051  usbmuxd    0.0 00:04.10 3      1      47      66      436K    216K    480K    60M     2422M
99006  iTunesHelper 0.0 00:01.23 2      1      55      78      728K    3124K   1124K   43M     2429M
84286  bash       0.0 00:00.11 1      0      20      24      224K    732K    484K    17M     2378M
84285  xterm      0.0 00:00.83 1      0      32      73      656K    872K    692K    9728K   2382M
55939- Microsoft Ex 0.3 21:58.97 10     3      360     954     16M     65M     46M     114M    1057M
54751  sleep      0.0 00:00.00 1      0      17      20      92K     212K    360K    9632K   2370M
54739  launchdadd 0.0 00:00.00 2      1      33      50      488K    220K    1736K   48M     2409M
54737  top        6.5 00:02.53 1/1    0      30      29      1416K   216K    2124K   17M     2378M
54719  automountd 0.0 00:00.02 7      1      53      64      860K    216K    2184K   53M     2413M
54701  ocspd      0.0 00:00.05 4      1      61      54      1268K   2644K   3132K   50M     2426M
54661  Grab       0.6 00:02.75 6      3      222+    389+    15M+    26M+    40M+    75M+    2556M+
54659  cookied    0.0 00:00.15 2      1      40      61      3316K   224K    4088K   42M     2411M
50678  mdownker  0.0 00:11.17 3      1      53      91      2464K   6148K   587M    44M     2434M
50610  xterm      0.0 00:00.13 1      0      32      73      280K    872K    532K    9700K   2382M
50609  emacs      0.0 00:06.70 1      0      20      35      52K     216K    88K     18M     2392M
```

## ■ 在Mac计算机上运行程序“top”命令 Running program “top” on Mac

- 系统有123个进程，5个是活跃状态 System has 123 processes, 5 of which are active
- 使用进程ID(PID)标识 Identified by Process ID (PID)



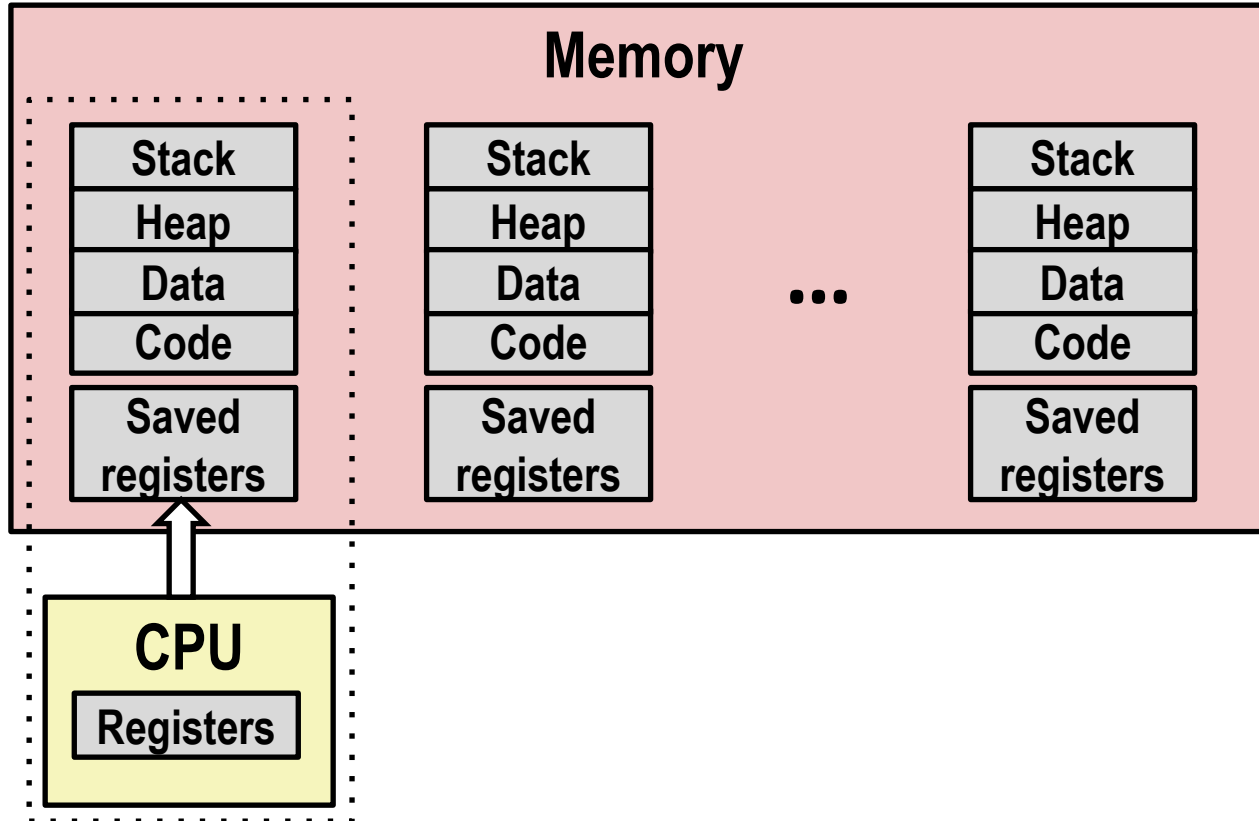
# 多进程的真像 Multiprocessing: The (Traditional) Reality



- **单个处理器并发执行多个进程** Single processor executes multiple processes concurrently
  - 进程交替执行（多任务） Process executions interleaved (multitasking)
  - 地址空间由虚拟内存系统管理 Address spaces managed by virtual memory system (later in course)
  - 非激活进程的寄存器值存储在内存中 Register values for nonexecuting processes saved in memory

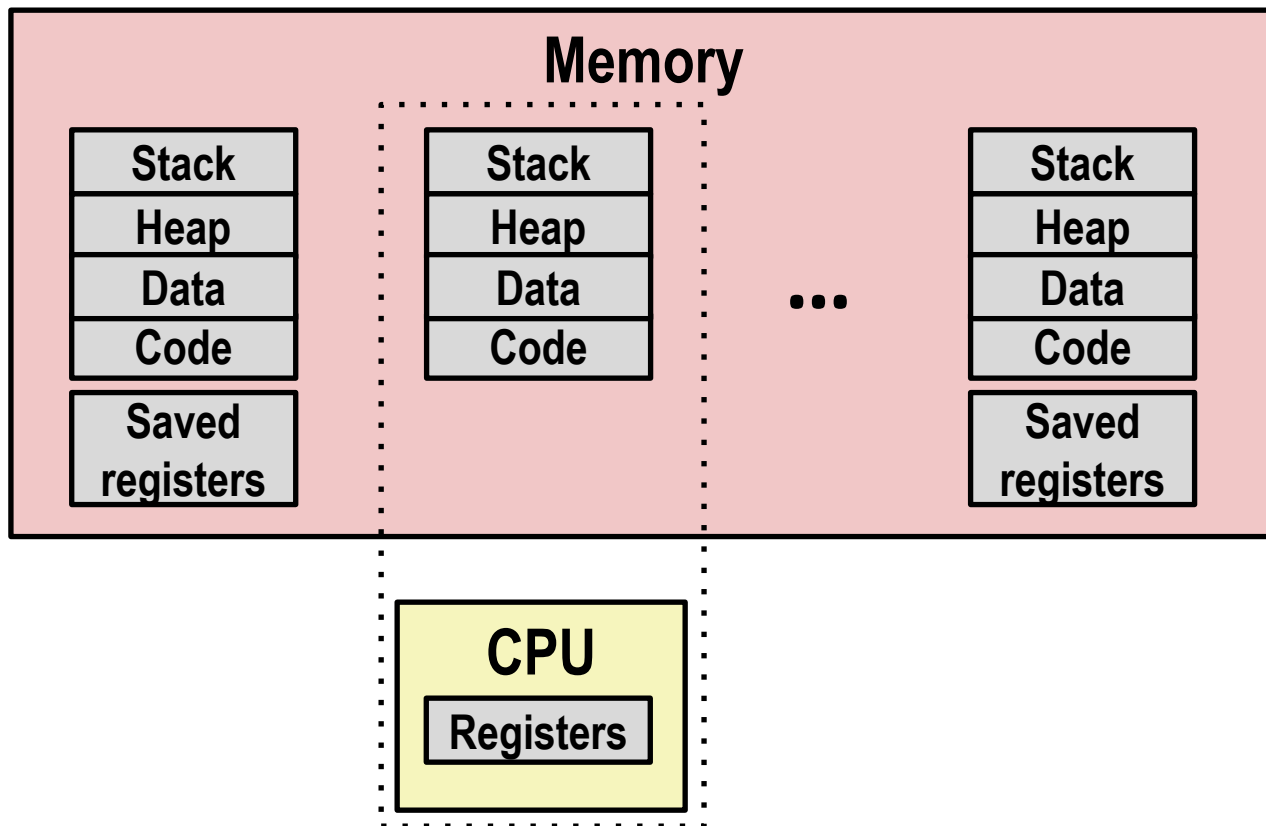


# 多进程真像 Multiprocessing: The (Traditional) Reality



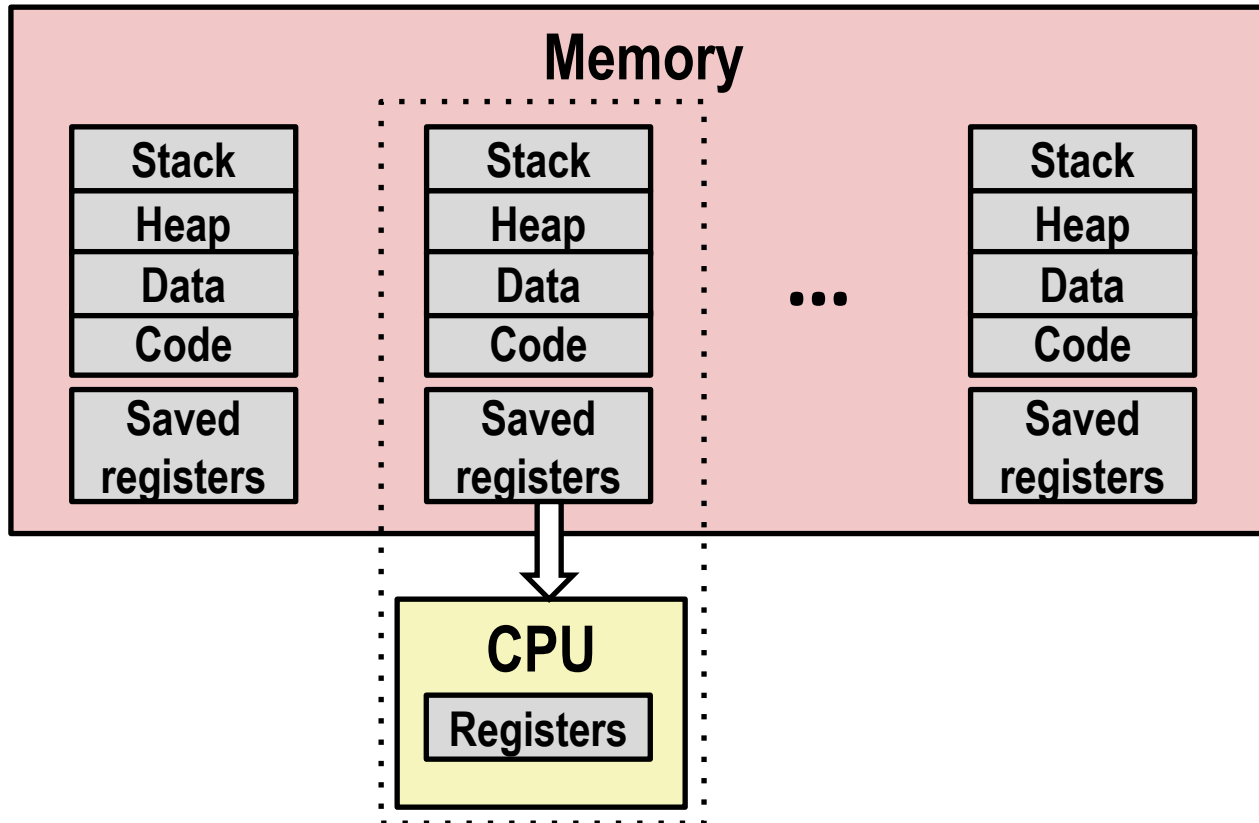
- 将当前寄存器存储在内存里 Save current registers in memory

# 多进程真像 Multiprocessing: The (Traditional) Reality



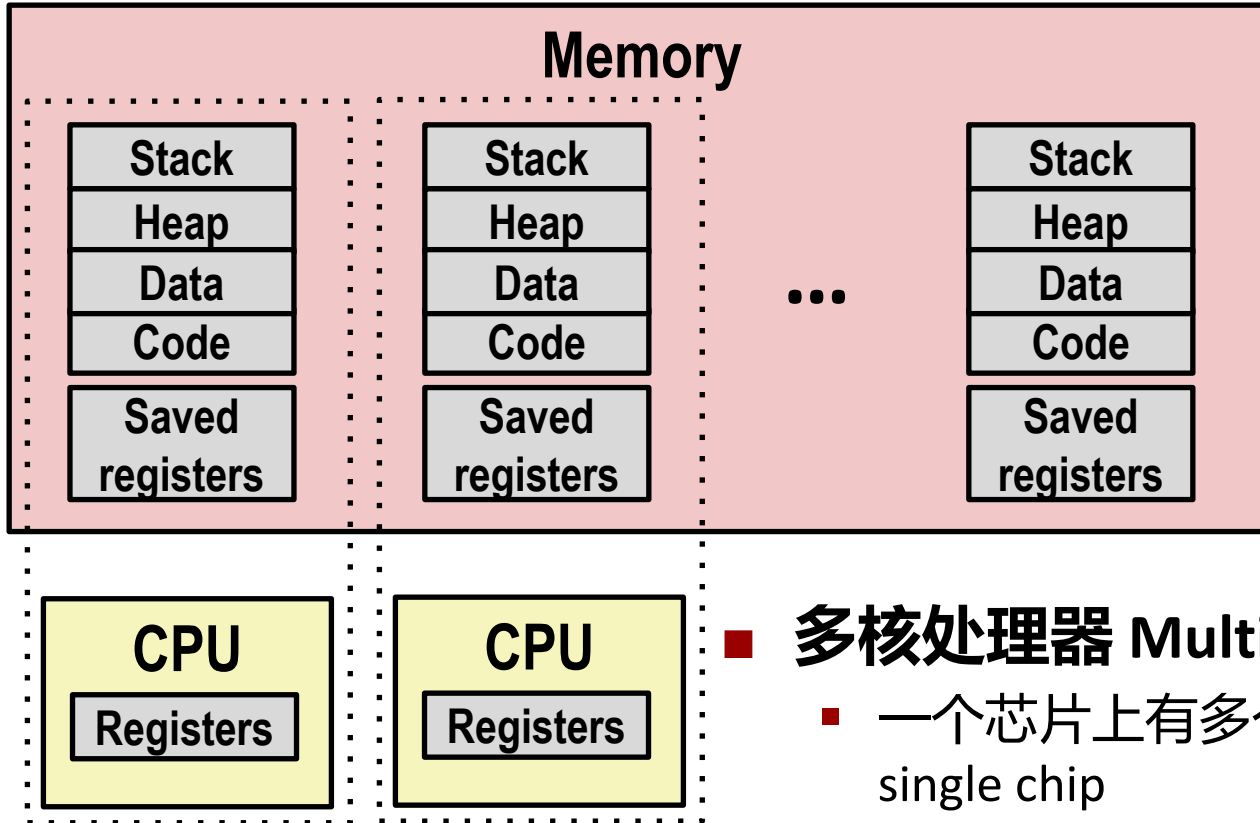
- 调度下一个进程执行 Schedule next process for execution

# 多进程真像 Multiprocessing: The (Traditional) Reality



- 加载保存的寄存器并切换地址空间（上下文切换） Load saved registers and switch address space (context switch)

# 多进程的真像 Multiprocessing: The (Modern) Reality



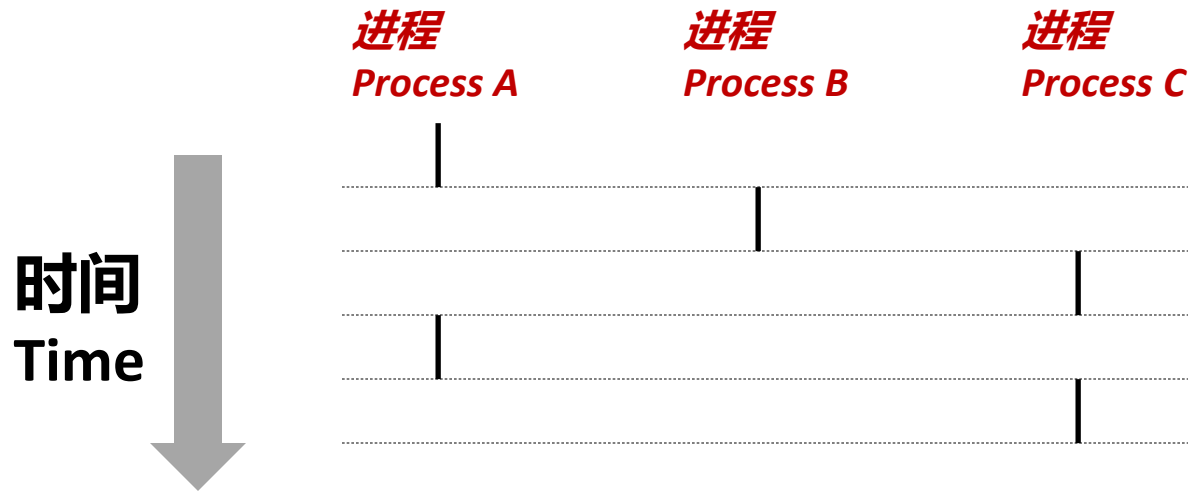
## ■ 多核处理器 Multicore processors

- 一个芯片上有多个CPU Multiple CPUs on single chip
- 共享主存储器（以及部分cache） Share main memory (and some of the caches)
- 每个可以执行一个独立进程 Each can execute a separate process
  - 由内核完成处理器到核心的调度 Scheduling of processors onto cores done by kernel



# 并发进程 Concurrent Processes

- 每个进程是一个逻辑控制流 Each process is a logical control flow.
- 两个进程**并发**运行如果在时间上重叠 Two processes *run concurrently* (are concurrent) if their flows overlap in time
- 否则是**顺序**执行 Otherwise, they are *sequential*
- 例如（运行在单核上） Examples (running on single core):
  - 并发：Concurrent: A & B, A & C
  - 顺序：Sequential: B & C

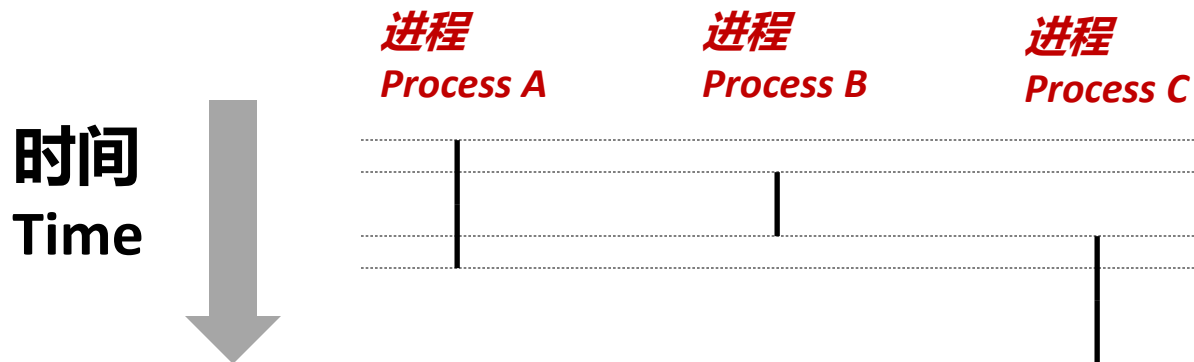




# 并发进程的用户视图

## User View of Concurrent Processes

- 并发进程的控制流在时间上是物理上不相交的 Control flows for concurrent processes are physically disjoint in time
- 然而，我们可以将并发进程视为彼此并行运行 However, we can think of concurrent processes as running in parallel with each other







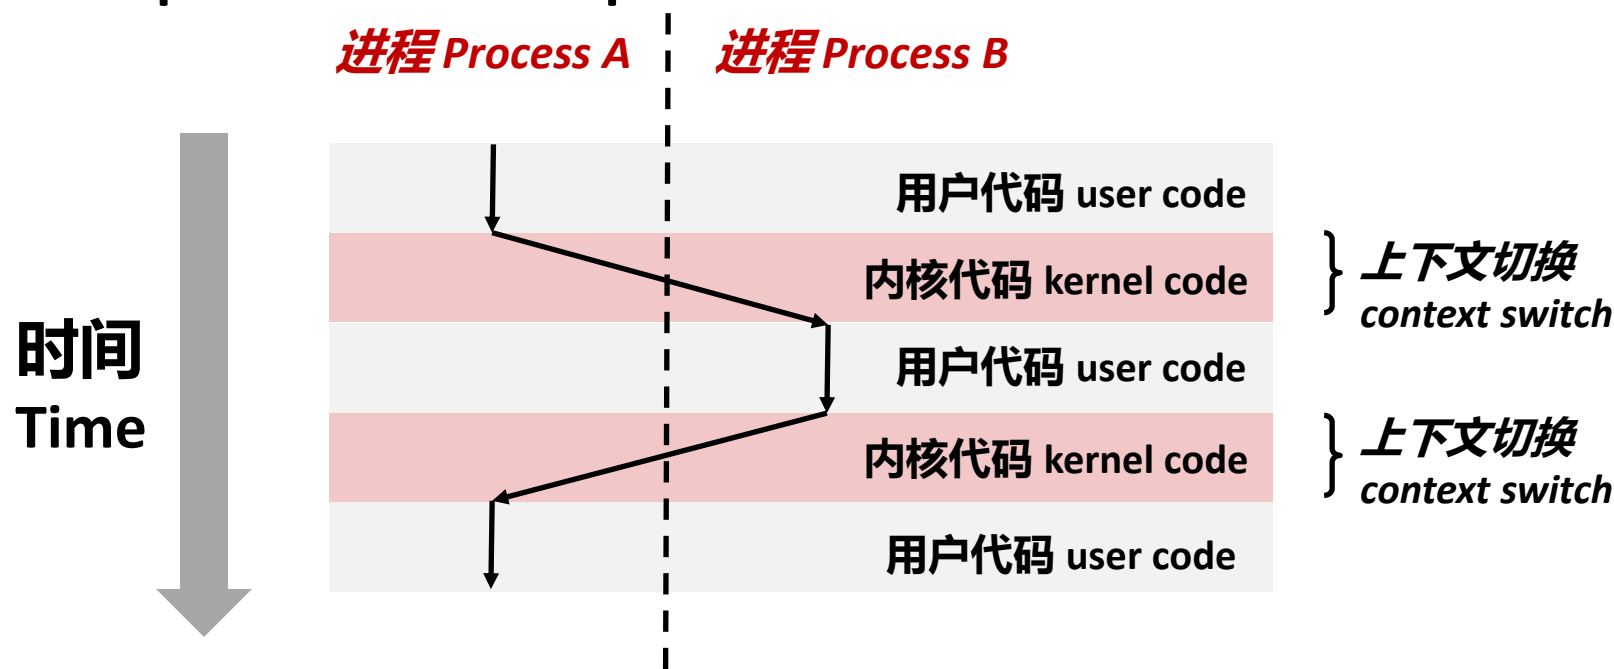
# 上下文切换 Context Switching

- 进程由称为**内核**的共享内存驻留操作系统代码块管理

Processes are managed by a shared chunk of memory-resident OS code called the **kernel**

- 重点：内核不是一个独立的进程，而是作为某些现存进程的一部分运行 Important: the kernel is not a separate process, but rather runs as part of some existing process.

- **上下文切换**使得控制流从一个进程切换到另一个进程 Control flow passes from one process to another via a **context switch**





# 内容提纲

- 异常控制流 Exceptional Control Flow
- 异常 Exceptions
- 进程 Processes
- 进程控制 Process Control

# 系统功能调用错误处理 System Call Error Handling



- **出错时，Linux系统函数返回-1并通过全局变量errno设置错误编号指明原因** On error, Linux system-level functions typically return -1 and set global variable `errno` to indicate cause.
- **硬性规定：Hard and fast rule:**
  - 你必须检查每个系统函数返回状态 You must check the return status of every system-level function
  - 返回值为void的函数除外 Only exception is the handful of functions that return `void`
- **例如 Example:**
- **#include <errno.h>**

```
if ((pid = fork()) < 0) {  
    fprintf(stderr, "fork error: %s\n", strerror(errno));  
    exit(0);  
}
```



# 错误报告函数 Error-reporting functions

- 使用错误报告函数可以简化一些工作 Can simplify somewhat using an *error-reporting function*:

```
void unix_error(char *msg) /* Unix-style error */  
{  
    fprintf(stderr, "%s: %s\n", msg, strerror(errno));  
    exit(0);  
}
```

```
if ((pid = fork()) < 0)  
    unix_error("fork error");
```

注意：退出时返回0  
Note: csapp.c exits with 0.

- 但是，必须考虑应用。当出现问题时退出并不总是合适的 But, must think about application. Not always appropriate to exit when something goes wrong.

# 错误处理包装器 Error-handling Wrappers



- 通过使用Stevens<sup>1</sup>风格的错误处理包装器，我们进一步简化了向您展示的代码： We simplify the code we present to you even further by using Stevens-style error-handling wrappers:

```
pid_t Fork(void)
{
    pid_t pid;

    if ((pid = fork()) < 0)
        unix_error("Fork error");
    return pid;
}
```

```
pid = Fork();
```

- 而不是您在实际应用程序中通常要做这件事情 NOT what you generally want to do in a real application

<sup>1</sup>例如，在“Unix网络编程：套接字网络API”<sup>1</sup>e.g., in “UNIX Network Programming: The sockets networking API” W. Richard Stevens



# 获得进程PID Obtaining Process IDs

- `pid_t getpid(void)`
  - 返回当前进程的PID Returns PID of current process
- `pid_t getppid(void)`
  - 返回父进程的PID Returns PID of parent process

# 创建和终止进程

## Creating and Terminating Processes



从程序员的角度，可以认为一个进程处于3种状态之一 From a programmer's perspective, we can think of a process as being in one of three states

### ■ 运行 Running

- 进程或者正在执行，或者等待被执行并最终由内核调度（即被选择执行） Process is either executing, or waiting to be executed and will eventually be *scheduled* (i.e., chosen to execute) by the kernel

### ■ 停止 Stopped

- 进程执行被挂起，直到被触发重新调度执行 Process execution is *suspended* and will not be scheduled until further notice (next lecture when we study signals)

### ■ 终止 Terminated

- 进程永远停止运行 Process is stopped permanently



# 进程终止 Terminating Processes




- 进程由于以下三个原因之一终止 Process becomes terminated for one of three reasons:
  - 收到默认动作是终止的信号 Receiving a signal whose default action is to terminate (next lecture)
  - 从main函数返回 Returning from the `main` routine
  - 调用exit函数 Calling the `exit` function
- `void exit(int status)`
  - 终止退出状态为status Terminates with an *exit status* of `status`
  - 规则：正常返回状态为0，出错为非0 Convention: normal return status is 0, nonzero on error
  - 另一种显式设置退出状态的方式是从main函数返回一个整数值  
Another way to explicitly set the exit status is to return an integer value from the main routine
- `exit`调用一次，但从不返回 `exit` is called **once** but **never** returns.

# 创建进程 Creating Processes



- **父进程通过调用 `fork` 创建一个新的运行子进程** *Parent process creates a new running child process by calling `fork`*
- **`int fork(void)`**
  - 返回0给子进程，子进程的PID给父进程 *Returns 0 to the child process, child's PID to parent process*
  - 子进程和父进程几乎是一样的 *Child is almost identical to parent:*
    - 子进程获得与父进程的虚拟地址空间同样的拷贝（但是是分开的） *Child get an identical (but separate) copy of the parent's virtual address space.*
    - 子进程获得与父进程打开文件描述符同样的拷贝 *Child gets identical copies of the parent's open file descriptors*
    - 子进程与父进程有不同的PID *Child has a different PID than the parent*
- **`fork` 很有意思（通常也令人费解），因为它调用一次，但返回两次** *`fork` is interesting (and often confusing) because it is called **once** but returns **twice***

# fork的概念视图 Conceptual View of fork



- **做完全的执行状态拷贝 Make complete copy of execution state**
  - 指定一个为父进程一个为子进程 Designate one as parent and one as child
  - 恢复父进程或子进程的执行 Resume execution of parent or child

# 重新审视fork函数



## The fork Function Revisited

- **虚拟存储器和内存映射解释了fork如何为每个进程提供私有的虚拟地址空间** VM and memory mapping explain how fork provides private address space for each process.
- **为了给新进程创建虚拟地址** To create virtual address for new process:
  - **创建与当前mm\_struct、vm\_area\_struct和页表精确一致的拷贝** Create exact copies of current mm\_struct, vm\_area\_struct, and page tables.
  - **设置两个进程对每个页具有只读权限** Flag each page in both processes as read-only
  - **设置两个进程对每个vm\_area\_struct都是私有COW** Flag each vm\_area\_struct in both processes as private COW
- **返回时，每个进程具有精确的虚拟内存拷贝** On return, each process has exact copy of virtual memory.
- **后续的写操作使用COW机制创建新页面** Subsequent writes create new pages using COW mechanism.



# fork举例 fork Example

```
int main(int argc, char** argv)
{
    pid_t pid;
    int x = 1;

    pid = Fork();
    if (pid == 0) { /* Child */
        printf("child : x=%d\n", ++x);
        return 0;
    }

    /* Parent */
    printf("parent: x=%d\n", --x);
    return 0;
}
```

*fork.c*

- 调用一次，返回两次 Call once, return twice
- 并发执行 Concurrent execution
  - 不能预测父进程和子进程的  
执行顺序 Can't predict  
execution order of parent  
and child

```
linux> ./fork
parent: x=0
child : x=2
```

```
linux> ./fork
child : x=2
parent: x=0
```

```
linux> ./fork
parent: x=0
child : x=2
```

```
linux> ./fork
parent: x=0
child : x=2
```

# fork举例 fork Example



```
int main(int argc, char** argv)
{
    pid_t pid;
    int x = 1;

    pid = Fork();
    if (pid == 0) { /* Child */
        printf("child : x=%d\n", ++x);
        return 0;
    }

    /* Parent */
    printf("parent: x=%d\n", --x);
    return 0;
}
```

## ■ 共享打开的文件 Shared open files

- 标准输出对父子进程是相同的 stdout is the same in both parent and child

```
linux> ./fork
parent: x=0
child : x=2
```

- 调用一次，返回两次 Call once, return twice
- 并发执行 Concurrent execution
  - 不能预测父进程和子进程的  
执行顺序 Can't predict  
execution order of parent  
and child
- 重复但是分开的地址空间 Duplicate but separate address space
  - x的值为1，当fork在父子进程返回 x has a value of 1 when fork returns in parent and child
  - 后续对x的改变是独立的 Subsequent changes to x are independent

# 使用进程图描述fork



## Modeling fork with Process Graphs

- **进程图**是一个有用的工具，它可以捕获并发程序中语句的偏序 *A process graph is a useful tool for capturing the partial ordering of statements in a concurrent program:*
  - 每个顶点都是语句的执行 Each vertex is the execution of a statement
  - $a \rightarrow b$ 表示a发生在b之前  $a \rightarrow b$  means a happens before b
  - 可以用变量的当前值标记边 Edges can be labeled with current value of variables
  - 可以用输出标记printf顶点 `printf` vertices can be labeled with output
  - 每个图都以一个没有输入边的顶点开始 Each graph begins with a vertex with no inedges
- **进程图的任何拓扑排序都对应于一种可行的全排序** *Any topological sort of the graph corresponds to a feasible total ordering.*
  - 所有边从左向右指向的顶点的全排序 Total ordering of vertices where all edges point from left to right




# 进程图举例 Process Graph Example

```
int main()
{
    pid_t pid;
    int x = 1;

    pid = Fork();
    if (pid == 0) { /* Child */
        printf("child : x=%d\n", ++x);
        exit(0);
    }

    /* Parent */
    printf("parent: x=%d\n", --x);
    exit(0);
}
```

*fork.c*









# 解释进程图 Interpreting Process Graphs

## ■ 原始图 Original graph:




## ■ 重新标记的图 Relabelled graph:




## 可行的全排序

Feasible total ordering:



## 可行还是不可行?

Feasible or Infeasible?



不可行: 不是一种拓扑排序

Infeasible: not a topological sort 45


# fork举例：两个连续的fork

## fork Example: Two consecutive forks



```
void fork2()
{
    printf("L0\n");
    fork();
    printf("L1\n");
    fork();
    printf("Bye\n");
}
```

*forks.c*



### 可能的输出

Feasible output:

L0  
L1  
Bye  
Bye  
L1  
Bye  
Bye

### 不可能的输出

Infeasible output:

L0  
Bye  
L1  
Bye  
L1  
Bye  
Bye




# fork举例：父类进程中的嵌套forks

## fork Example: Nested forks in parent

```
void fork4()
{
    printf("L0\n");
    if (fork() != 0) {
        printf("L1\n");
        if (fork() != 0) {
            printf("L2\n");
        }
    }
    printf("Bye\n");
}
```

*forks.c*



### 可能的输出

Feasible output:

L0  
L1  
Bye  
Bye  
L2  
Bye

### 不可能的输出

Infeasible output:

L0  
Bye  
L1  
Bye  
Bye  
L2


# fork举例：子进程中的嵌套forks

## fork Example: Nested forks in children



```
void fork5()
{
    printf("L0\n");
    if (fork() == 0) {
        printf("L1\n");
        if (fork() == 0) {
            printf("L2\n");
        }
    }
    printf("Bye\n");
}
```

*forks.c*



### 可能的输出

Feasible output:

L0  
Bye  
L1  
L2  
Bye  
Bye

### 不可能的输出

Infeasible output:

L0  
Bye  
L1  
Bye  
Bye  
L2

# 回收子进程 Reaping Child Processes



## ■ 思想 Idea

- 进程终止后仍然消耗系统资源 When process terminates, it still consumes system resources
  - 例如：退出状态，各种OS表格 Examples: Exit status, various OS tables
- 称为“僵尸”进程 Called a “zombie”
  - 活着的尸体，半生半死 Living corpse, half alive and half dead

## ■ 回收 Reaping

- 父类进程对终止的子进程操作（使用wait或waitpid） Performed by parent on terminated child (using `wait` or `waitpid`)
- 父类进程持有退出状态信息 Parent is given exit status information
- 内核随后删掉僵尸子进程 Kernel then deletes zombie child process

# 回收子进程 Reaping Child Processes



- 如果父类进程没有回收会怎么样？ What if parent doesn't reap?
  - 如果任何父类进程终止没有回收子进程，则该孤儿子进程由init进程（pid==1）回收 If any parent terminates without reaping a child, then the orphaned child will be reaped by **init** process (pid == 1)
    - 除非ppid==1，此时需要重启 Unless ppid == 1! Then need to reboot...
  - 所以只需要显式回收长时间运行的进程 So, only need explicit reaping in long-running processes
    - 例如外壳程序和服务器程序 e.g., shells and servers

# 僵尸举例

## Zombie Example



```
void fork7() {  
    if (fork() == 0) {  
        /* Child */  
        printf("Terminating Child, PID = %d\n", getpid());  
        exit(0);  
    } else {  
        printf("Running Parent, PID = %d\n", getpid());  
        while (1)  
            ; /* Infinite loop */  
    }  
}
```

```
linux> ./forks 7 &  
[1] 6639
```

```
Running Parent, PID = 6639
```

```
Terminating Child, PID = 6640
```

```
linux> ps
```

PID	TTY	TIME	CMD
6585	ttyp9	00:00:00	tcsh
6639	ttyp9	00:00:03	forks
6640	ttyp9	00:00:00	forks <defunct>
6641	ttyp9	00:00:00	ps

```
linux> kill 6639
```

```
[1] Terminated
```

```
linux> ps
```

PID	TTY	TIME	CMD
6585	ttyp9	00:00:00	tcsh
6642	ttyp9	00:00:00	ps

■ ps显示子进程为“defunct”  
(即僵尸) ps shows child  
process as “defunct” (i.e., a  
zombie)

■ 杀死父进程允许子进程由init  
进程回收 Killing parent allows  
child to be reaped by **init**

# 非终止子进程举例

## Non-terminating Child Example

```
void fork8()
{
    if (fork() == 0) {
        /* Child */
        printf("Running Child, PID = %d\n",
               getpid());
        while (1)
            ; /* Infinite loop */
    } else {
        printf("Terminating Parent, PID = %d\n",
               getpid());
        exit(0);
    }
}
```

```
linux> ./forks 8
Terminating Parent, PID = 6675
Running Child, PID = 6676
linux> ps
  PID TTY          TIME CMD
 6585 tttyp9      00:00:00 tcsh
 6676 tttyp9      00:00:06 forks
 6677 tttyp9      00:00:00 ps
linux> kill 6676
linux> ps
  PID TTY          TIME CMD
 6585 tttyp9      00:00:00 tcsh
 6678 tttyp9      00:00:00 ps
```

- 子进程仍然活着，尽管父进程已经终止 Child process still active even though parent has terminated
- 必须显式杀死子进程，否则子进程将会永远一直在运行 Must kill child explicitly, or else will keep running indefinitely






# wait: 与子进程同步

## wait: Synchronizing with Children

- 父进程通过调用wait函数回收子进程 Parent reaps a child by calling the wait function
- `int wait(int *child_status)`
  - 挂起当前进程直到其子进程之一终止 Suspend current process until one of its children terminates
  - 用syscall实现 Implemented as syscall



而且，潜在地其它用户进程，包括父进程的子进程  
And, potentially other user processes, including a child of parent



# wait: 与子进程同步

## wait: Synchronizing with Children

- 父进程通过调用wait函数回收子进程 Parent reaps a child by calling the wait function
- `int wait(int *child_status)`
  - 挂起当前进程直到其子进程之一终止 Suspends current process until one of its children terminates
  - 返回值是终止子进程的PID Return value is the `pid` of the child process that terminated
  - 如果`child_status`不为空，那么它指向的整数将会设置为一个值，以指示子进程终止的原因和退出状态： If `child_status != NULL`, then the integer it points to will be set to a value that indicates reason the child terminated and the exit status:
    - 使用wait.h中宏定义进行检查 Checked using macros defined in `wait.h`
      - `WIFEXITED`, `WEXITSTATUS`, `WIFSIGNALED`, `WTERMSIG`, `WIFSTOPPED`, `WSTOPSIG`, `WIFCONTINUED`
      - 参见教材了解详情 See textbook for details (P553、517)




# wait: 与子进程同步

## wait: Synchronizing with Children

```
void fork9() {  
    int child_status;  
  
    if (fork() == 0) {  
        printf("HC: hello from child\n");  
        exit(0);  
    } else {  
        printf("HP: hello from parent\n");  
        wait(&child_status);  
        printf("CT: child has terminated\n");  
    }  
    printf("Bye\n");  
}
```

*forks.c*



### 可能的输出

Feasible output(s):

HC  
HP  
CT  
Bye

### 不可能的输出

Infeasible output:

HP  
CT  
Bye  
HC

# 另一个wait的例子



## Another wait Example

- 如果多个子进程终止，将会以任意顺序进行 If multiple children completed, will take in arbitrary order
- 可以使用宏WIFEXITED和WEXITSTATUS获取有关退出状态的信息 Can use macros WIFEXITED and WEXITSTATUS to get information about exit status

```
void fork10() {  
    pid_t pid[N];  
    int i, child_status;  
  
    for (i = 0; i < N; i++)  
        if ((pid[i] = fork()) == 0) {  
            exit(100+i); /* Child */  
        }  
    for (i = 0; i < N; i++) { /* Parent */  
        pid_t wpid = wait(&child_status);  
        if (WIFEXITED(child_status))  
            printf("Child %d terminated with exit status %d\n",  
                wpid, WEXITSTATUS(child_status));  
        else  
            printf("Child %d terminate abnormally\n", wpid);  
    }  
}
```

*forks.c*



# waitpid: 等待特定进程

## waitpid: Waiting for a Specific Process

- `pid_t waitpid(pid_t pid, int &status, int options)`
  - 挂起当前进程直到指定进程终止 Suspends current process until specific process terminates
  - 各种选项（参见教材） Various options (see textbook)

```
void fork11() {  
    pid_t pid[N];  
    int i;  
    int child_status;  
  
    for (i = 0; i < N; i++)  
        if ((pid[i] = fork()) == 0)  
            exit(100+i); /* Child */  
    for (i = N-1; i >= 0; i--) {  
        pid_t wpid = waitpid(pid[i], &child_status, 0);  
        if (WIFEXITED(child_status))  
            printf("Child %d terminated with exit status %d\n",  
                wpid, WEXITSTATUS(child_status));  
        else  
            printf("Child %d terminate abnormally\n", wpid);  
    }  
}
```

*forks.c*



# execve: 加载运行程序


## execve: Loading and Running Programs

- `int execve(char *filename, char *argv[], char *envp[])`
- **在当前进程加载和运行** Loads and runs in the current process:
  - 可执行文件文件名filename Executable file **filename**
    - 目标代码文件或者以“#! 解释器”开始的脚本文件 Can be object file or script file beginning with `#!interpreter` (e.g., `#!/bin/bash`)
  - ...带有参数列表argv ...with argument list **argv**
    - 按照约定第一个参数为文件名 By convention **argv[0]==filename**
  - ...带有环境变量列表envp ...and environment variable list **envp**
    - “名字=值”串 “name=value” strings (e.g., `USER=droh`)
    - `getenv`, `putenv`, `putenv`
- **覆盖代码、数据和堆栈** Overwrites code, data, and stack
  - 保持PID、打开文件和信号上下文 Retains PID, open files and signal context
- **调用一次而且从不返回** Called **once** and **never** returns
  - ...除非如果有错误才返回调用程序 ...except if there is an error
    - 找不到filename

# Execve举例 execve Example



- 使用当前环境在子进程中执行 Execute `"/bin/ls -lt /usr/include"` in child process using current environment:

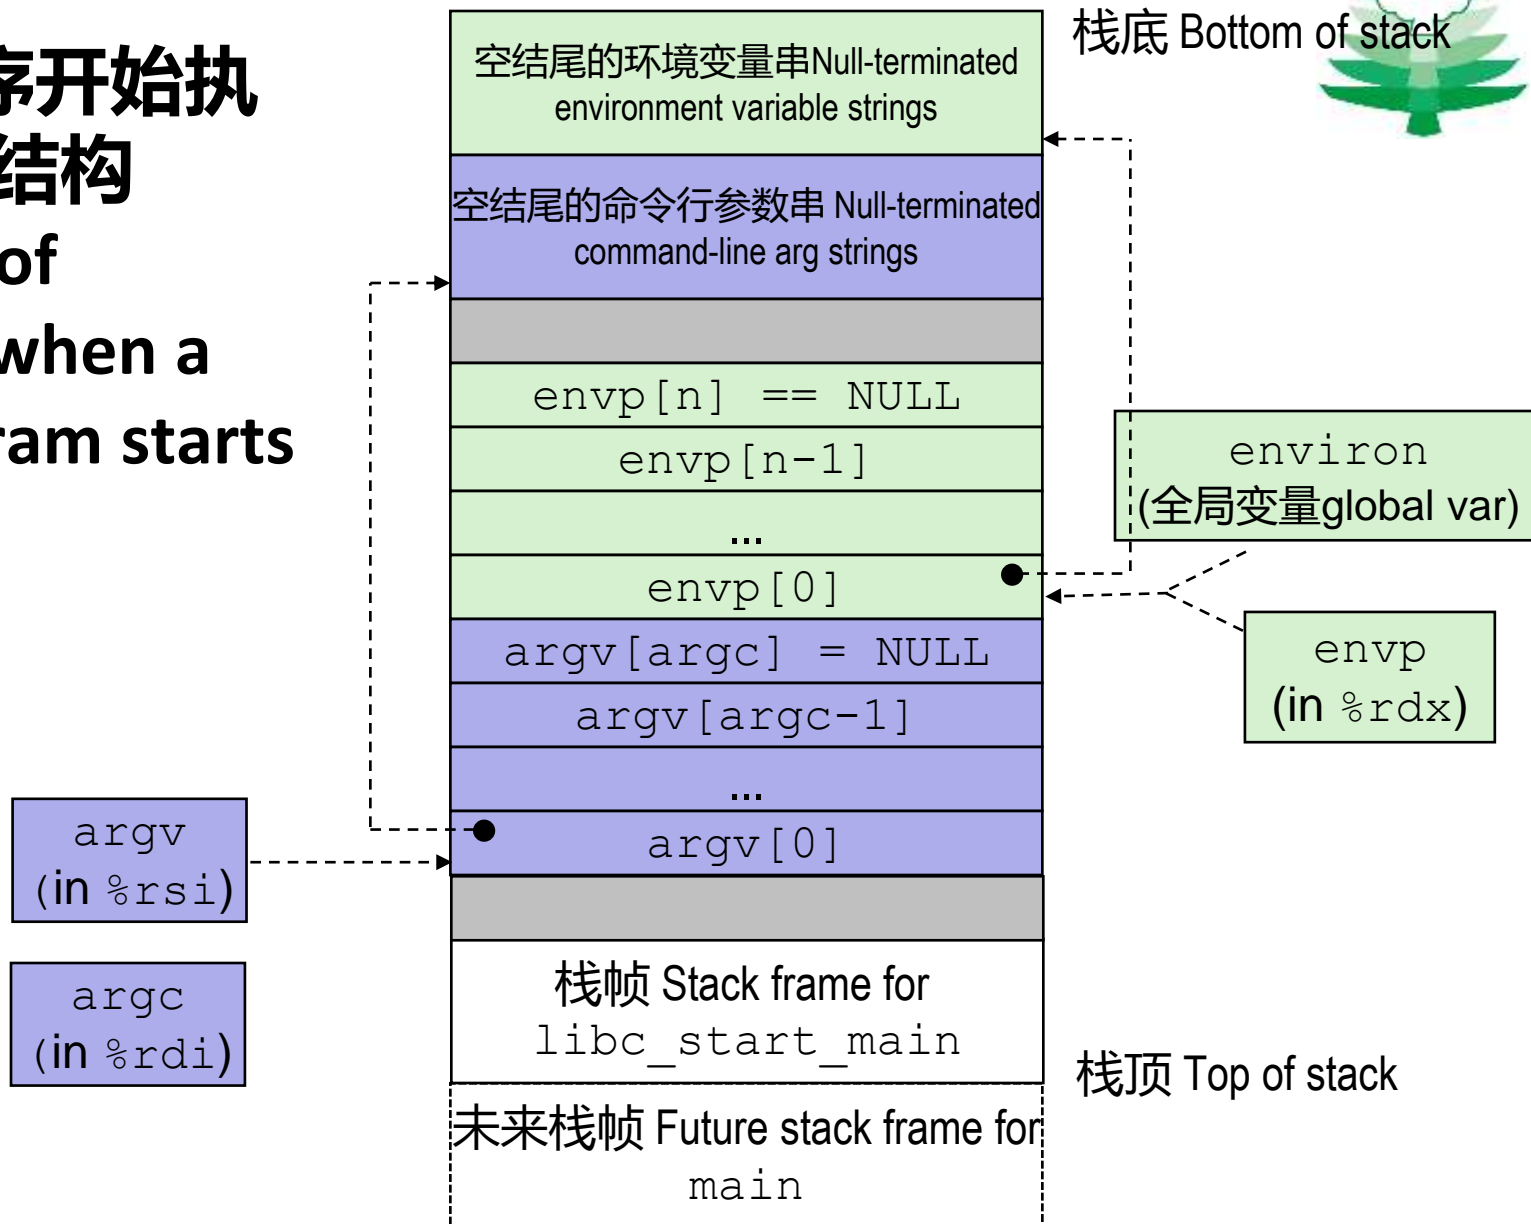


```
if ((pid = Fork()) == 0) {    /* Child runs program */
    if (execve(myargv[0], myargv, environ) < 0) {
        printf("%s: Command not found.\n", myargv[0]);
        exit(1);
    }
}
```



# 一个新程序开始执行时的栈结构

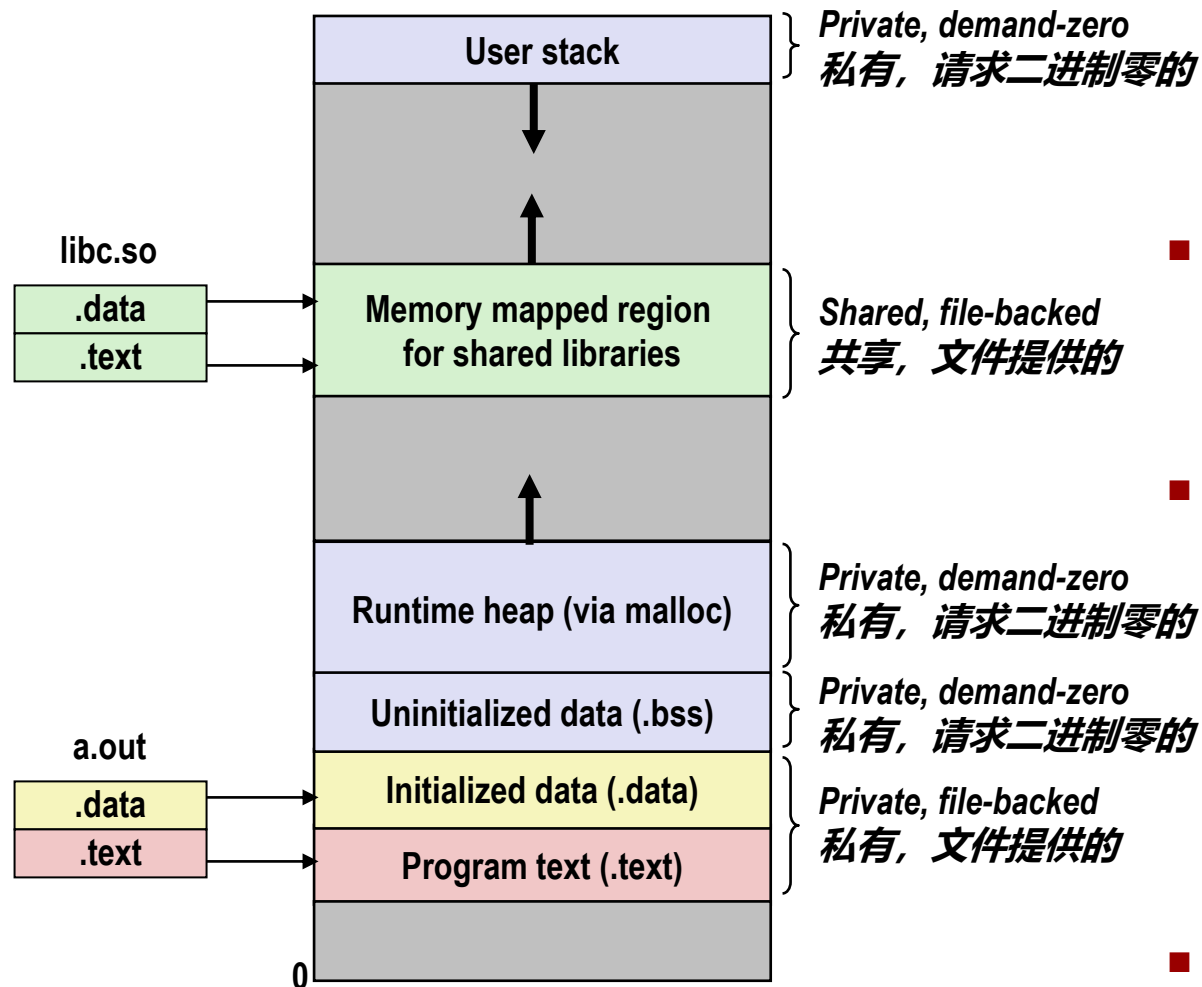
## Structure of the stack when a new program starts





# 重新审视execve函数

## The execve Function Revisited



- 要使用`execve`在当前进程加载和运行一个新程序`a.out` To load and run a new program `a.out` in the current process using `execve`:
- 释放老区域的`vm_area_struct`和页表 Free `vm_area_struct`'s and page tables for old areas
- 为新区域创建`vm_area_struct`和页表 Create `vm_area_struct`'s and page tables for new areas
  - 程序和初始化后的数据由目标文件提供 Programs and initialized data backed by object files.
  - `.bss`和栈由匿名文件提供 `.bss` and stack backed by anonymous files.
- 设置PC为`.text`中的入口点 Set PC to entry point in `.text`
  - Linux将陷入需要的代码和数据页 Linux will fault in code and data pages as needed.



# 总结 Summary

## ■ 异常 Exceptions

- 需要非标准控制流的事件 Events that require nonstandard control flow
- 由外部（中断）或内部（陷阱和故障）产生 Generated externally (interrupts) or internally (traps and faults)

## ■ 进程 Processes

- 任意时刻，系统有多个活动进程 At any given time, system has multiple active processes
- 尽管在单核上每个时刻只能执行一个进程 Only one can execute at a time on a single core, though
- 每个进程看起来独占处理器和私有内存空间 Each process appears to have total control of processor + private memory space



# 总结（续） Summary (cont.)

- **生成新进程 Spawning processes**
  - 调用fork Call `fork`
  - 一次调用，两次返回 One call, two returns
- **结束进程 Process completion**
  - 调用exit Call `exit`
  - 一次调用，不返回 One call, no return
- **回收和等待进程 Reaping and waiting for processes**
  - 调用wait或waitpid Call `wait` or `waitpid`
- **加载运行程序 Loading and running programs**
  - 调用execve（或变种） Call `execve` (or variant)
  - 一次调用，（正常）不返回 One call, (normally) no return

# 使fork更不确定

## Making `fork` More Nondeterministic



### ■ 问题 Problem

- Linux调度器不会产生很多运行间差异 Linux scheduler does not create much run-to-run variance
- 在非确定性程序中隐藏潜在的竞争条件 Hides potential race conditions in nondeterministic programs
  - 例如，`fork`是先返回到子进程，还是返回到父进程？ E.g., does `fork` return to child first, or to parent?

### ■ 解决方案 Solution

- 创建库例程的自定义版本，沿不同分支插入随机延迟 Create custom version of library routine that inserts random delays along different branches
  - 例如，`fork`父进程和子进程 E.g., for parent and child in `fork`
- 使用运行时库打桩使程序使用特殊版本的库代码 Use runtime interpositioning to have program use special version of library code

# 延迟变化的fork Variable delay fork



```
/* fork wrapper function */
pid_t fork(void) {
    initialize();
    int parent_delay = choose_delay();
    int child_delay = choose_delay();
    pid_t parent_pid = getpid();
    pid_t child_pid_or_zero = real_fork();
    if (child_pid_or_zero > 0) {
        /* Parent */
        if (verbose) {
            printf(
"Fork.  Child pid=%d, delay = %dms.  Parent pid=%d, delay = %dms\n",
                child_pid_or_zero, child_delay,
                parent_pid, parent_delay);
            fflush(stdout);
        }
        ms_sleep(parent_delay);
    } else {
        /* Child */
        ms_sleep(child_delay);
    }
    return child_pid_or_zero;
}
```



# 第8章 异常控制流 第二讲

## 信号和非本地跳转 Signals and Nonlocal Jumps

100076202: 计算机系统导论



任课教师:

宿红毅 张艳 黎有琦 颜珂

原作者:

Randal E. Bryant and David R. O'Hallaron

Carnegie  
Mellon  
University



# 异常控制流存在系统每个层次

## ECF Exists at All Levels of a System

### ■ 异常 Exceptions

- 硬件和操作系统内核软件
- Hardware and operating system kernel software

### ■ 进程上下文切换 Process Context Switch

- 硬件时钟和内核软件
- Hardware timer and kernel software

### ■ 信号 Signals

- 内核软件和应用软件
- Kernel software and application software

### ■ 非局部跳转 Nonlocal jumps

- 应用代码 Application code

Previous Lecture  
前面的课

This Lecture  
本次课

教材和补充幻灯片  
Textbook and  
supplemental slides

# (部分) 分类

## (partial) Taxonomy

内核处理 Handled in kernel

用户进程处理 Handled in user process

异常控制流ECF

```
graph TD; ECF[异常控制流ECF] --> Async[异步Asynchronous]; ECF --> Sync[同步 Synchronous]; Async --> Interrupts[中断Interrupts]; Async --> Signals[信号Signals]; Sync --> Traps[陷阱 Traps]; Sync --> Faults[故障 Faults]; Sync --> Aborts[终止Aborts];
```

异步Asynchronous

同步 Synchronous

中断Interrupts

信号Signals

陷阱 Traps

故障 Faults

终止Aborts



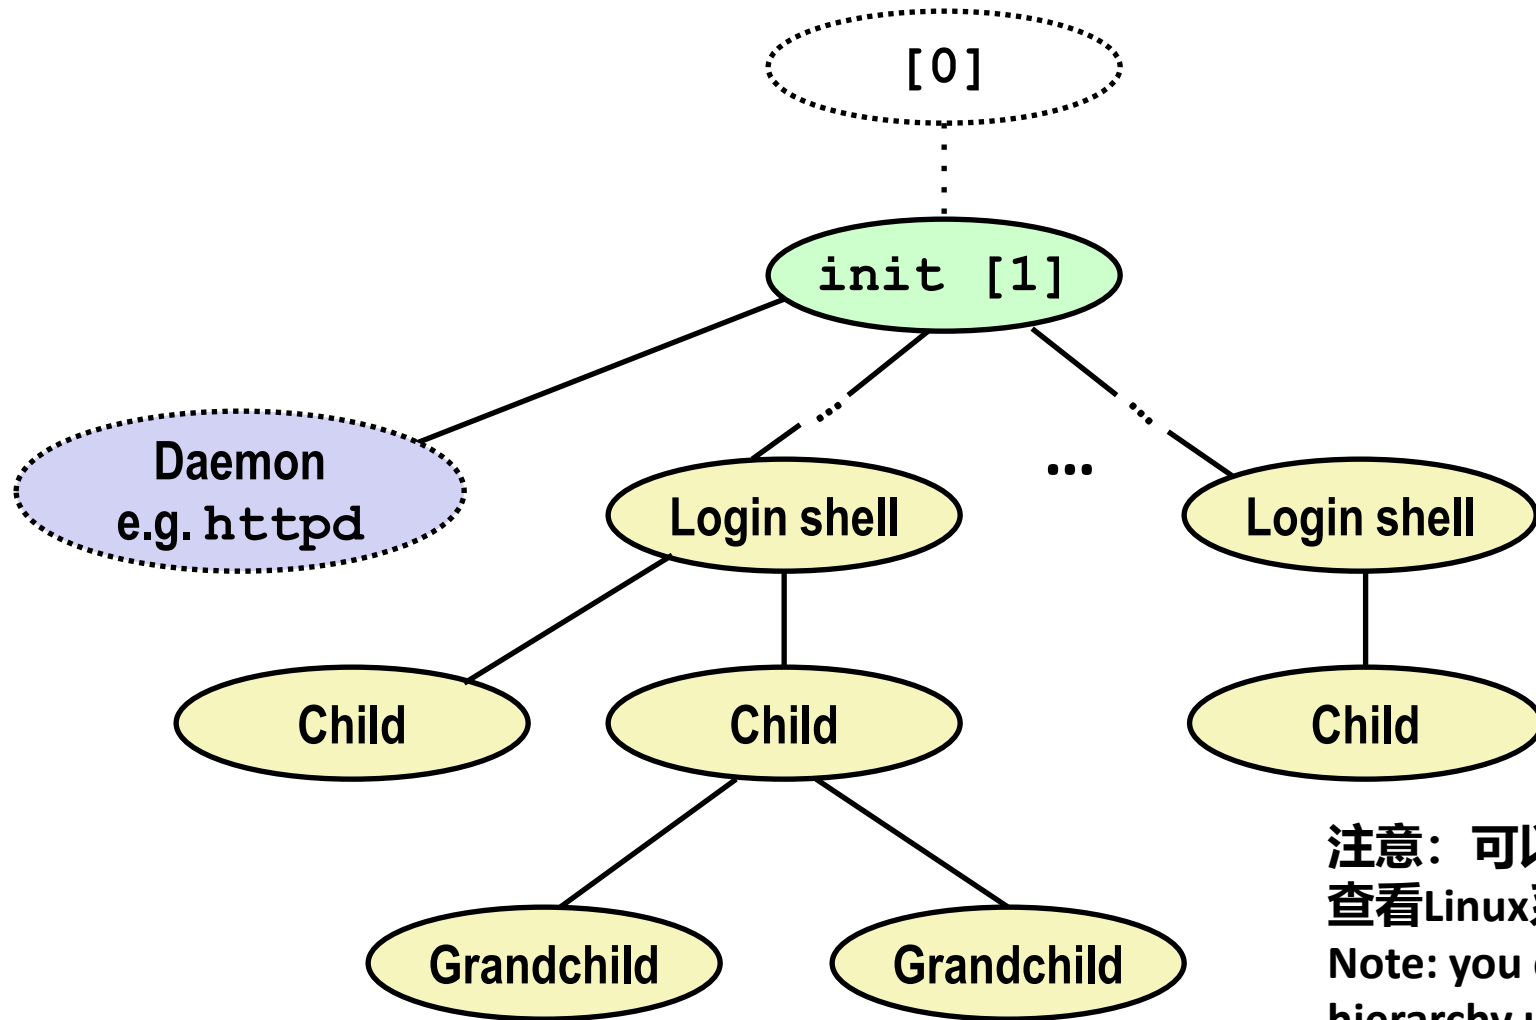


# 议题

- 外壳 Shells
- 信号 Signals
- 非局部跳转 Nonlocal jumps



# Linux进程树 Linux Process Hierarchy



注意：可以用`ps tree`命令  
查看Linux系统的进程树  
Note: you can view the  
hierarchy using the Linux  
`ps tree` command



# Shell程序 Shell Programs

- Shell是按照用户要求运行程序的应用程序 A *shell* is an application program that runs programs on behalf of the user
  - `sh` 最早的 Original Unix shell (Stephen Bourne, AT&T Bell Labs, 1977)
  - `csh/tcsh` BSD Unix C shell
  - `bash` 默认的 “Bourne-Again” Shell (default Linux shell)
- 简单shell Simple shell
  - 教材p524页处描述 Described in the textbook, starting at p.524
  - 一个非常基础的shell实现 Implementation of a very elementary shell
  - 目的 Purpose
    - 理解当输入了命令后究竟发生了什么事情 Understand what happens when you type commands
    - 理解进程控制操作的使用和操作 Understand use and operation of process control operations



# 简单shell示例 Simple Shell Example

```
linux> ./shellex
> /bin/ls -l csapp.c 必须给出程序的全路径名 Must give full pathnames for programs
-rw-r--r-- 1 bryant users 23053 Jun 15 2015 csapp.c
> /bin/ps
  PID TTY          TIME CMD
 31542 pts/2        00:00:01 tcsh
 32017 pts/2        00:00:00 shellex
 32019 pts/2        00:00:00 ps
> /bin/sleep 10 & 后台运行程序 Run program in background
32031 /bin/sleep 10 &
> /bin/ps
  PID TTY          TIME CMD
 31542 pts/2        00:00:01 tcsh
 32024 pts/2        00:00:00 emacs
 32030 pts/2        00:00:00 shellex
 32031 pts/2        00:00:00 sleep Sleep正在后台运行
 32033 pts/2        00:00:00 ps Sleep is running
> quit in background
```

# 简单shell实现

## Simple Shell Implementation



### ■ 基本循环 Basic loop

- 从命令行读一行 Read line from command line
- 执行请求的操作 Execute the requested operation
  - 内置命令（仅实现一个命令是**quit**） Built-in command (only one implemented is **quit**)
  - 从文件加载和执行程序 Load and execute program from file

```
int main(int argc, char** argv)
{
    char cmdline[MAXLINE]; /* command line */

    while (1) {
        /* read */
        printf("> ");
        fgets(cmdline, MAXLINE, stdin);
        if (feof(stdin))
            exit(0);

        /* evaluate */
        eval(cmdline);
    }
    ...
}
```

*shellex.c*

*执行的过程就是一系列读/求值的步骤 Execution is a sequence of read/evaluate steps*

# 简单的Shell eval函数 Simple Shell eval Function



```
void eval(char *cmdline)
{
    char *argv[MAXARGS]; /* Argument list execve() */
    char buf[MAXLINE]; /* Holds modified command line */
    int bg; /* Should the job run in bg or fg? */
    pid_t pid; /* Process id */

    strcpy(buf, cmdline);
    bg = parseline(buf, argv);
    if (argv[0] == NULL)
        return; /* Ignore empty lines */

    if (!builtin_command(argv)) {
        if ((pid = Fork()) == 0) { /* Child runs user job */
            if (execve(argv[0], argv, environ) < 0) {
                printf("%s: Command not found.\n", argv[0]);
                exit(0);
            }
        }

        /* Parent waits for foreground job to terminate */
        if (!bg) {
            int status;
            if (waitpid(pid, &status, 0) < 0)
                unix_error("waitfg: waitpid error");
        }
        else
            printf("%d %s", pid, cmdline);
    }
    return;
}
```

# 简单的Shell eval函数 Simple Shell eval Function



```
void eval(char *cmdline)
{
    char *argv[MAXARGS]; /* Argument list execve() */
    char buf[MAXLINE];    /* Holds modified command line */
    int bg;               /* Should the job run in bg or fg? */
    pid_t pid;           /* Process id */

    strcpy(buf, cmdline);
    bg = parseline(buf, argv);
}
```

**Parseline函数将buf解析成argv并返回是否输入行以&结尾**  
parseline will parse 'buf' into 'argv' and return whether or not input line ended in '&'

# 简单的Shell eval函数 Simple Shell eval Function



```
void eval(char *cmdline)
{
    char *argv[MAXARGS]; /* Argument list execve() */
    char buf[MAXLINE];    /* Holds modified command line */
    int bg;               /* Should the job run in bg or fg? */
    pid_t pid;            /* Process id */

    strcpy(buf, cmdline);
    bg = parseline(buf, argv);
    if (argv[0] == NULL)
        return; /* Ignore empty lines */
```

忽略空行  
Ignore empty lines.



# 简单的Shell eval函数 Simple Shell eval Function



```
void eval(char *cmdline)
{
    char *argv[MAXARGS]; /* Argument list execve() */
    char buf[MAXLINE];    /* Holds modified command line */
    int bg;               /* Should the job run in bg or fg? */
    pid_t pid;            /* Process id */

    strcpy(buf, cmdline);
    bg = parseline(buf, argv);
    if (argv[0] == NULL)
        return; /* Ignore empty lines */

    if (!builtin_command(argv)) {
```

如果是“内置”命令，那么在这个程序此处处理它。否则创建进程(fork)/执行(exec) 在argv[0]中指定的程序  
If it is a 'built in' command, then handle it here in this program. Otherwise fork/exec the program specified in argv[0]

# 简单的Shell eval函数 Simple Shell eval Function



```
void eval(char *cmdline)
{
    char *argv[MAXARGS]; /* Argument list execve() */
    char buf[MAXLINE];    /* Holds modified command line */
    int bg;               /* Should the job run in bg or fg? */
    pid_t pid;            /* Process id */

    strcpy(buf, cmdline);
    bg = parseline(buf, argv);
    if (argv[0] == NULL)
        return; /* Ignore empty lines */

    if (!builtin_command(argv)) {
        if ((pid = fork()) == 0) { /* Child runs user job */
```

创建子进程/Create child

# 简单的Shell eval函数 Simple Shell eval Function



```
void eval(char *cmdline)
{
    char *argv[MAXARGS]; /* Argument list execve() */
    char buf[MAXLINE];    /* Holds modified command line */
    int bg;               /* Should the job run in bg or fg? */
    pid_t pid;            /* Process id */

    strcpy(buf, cmdline);
    bg = parseline(buf, argv);
    if (argv[0] == NULL)
        return; /* Ignore empty lines */

    if (!builtin_command(argv)) {
        if ((pid = fork()) == 0) { /* Child runs user job */
            if (execve(argv[0], argv, environ) < 0) {
                printf("%s: Command not found.\n", argv[0]);
                exit(0);
            }
        }
    }
}
```

启动argv[0].

记住execve仅在出错时返回

Start **argv[0]**.

Remember **execve** only returns on error.

# 简单的Shell eval函数 Simple Shell eval Function



```
void eval(char *cmdline)
{
    char *argv[MAXARGS]; /* Argument list execve() */
    char buf[MAXLINE];    /* Holds modified command line */
    int bg;               /* Should the job run in bg or fg? */
    pid_t pid;            /* Process id */

    strcpy(buf, cmdline);
    bg = parseline(buf, argv);
    if (argv[0] == NULL)
        return; /* Ignore empty lines */

    if (!builtin_command(argv)) {
        if ((pid = fork()) == 0) { /* Child runs user job */
            if (execve(argv[0], argv, environ) < 0) {
                printf("%s: Command not found.\n", argv[0]);
                exit(0);
            }
        }

        /* Parent waits for foreground job to terminate */
        if (!bg) {
            int status;
            if (waitpid(pid, &status, 0) < 0)
                unix_error("waitfg: waitpid error");
        }
    }
}
```

如果子进程在前台运行，等待直到子进程完成  
If running child in foreground, wait until it is done.

# 简单的Shell eval函数 Simple Shell eval Function



```
void eval(char *cmdline)
{
    char *argv[MAXARGS]; /* Argument list execve() */
    char buf[MAXLINE];    /* Holds modified command line */
    int bg;               /* Should the job run in bg or fg? */
    pid_t pid;            /* Process id */

    strcpy(buf, cmdline);
    bg = parseline(buf, argv);
    if (argv[0] == NULL)
        return; /* Ignore empty lines */

    if (!builtin_command(argv)) {
        if ((pid = fork()) == 0) { /* Child runs user job */
            if (execve(argv[0], argv, environ) < 0) {
                printf("%s: Command not found.\n", argv[0]);
                exit(0);
            }
        }

        /* Parent waits for foreground job to terminate */
        if (!bg) {
            int status;
            if (waitpid(pid, &status, 0) < 0)
                unix_error("waitfg: waitpid error");
        }
        else
            printf("%d %s", pid, cmdline);
    }
    return;
}
```

如果子进程在后台运行，打印pid并继续做其它事情

If running child in background, print pid and continue doing other stuff.

# 简单的Shell eval函数 Simple Shell eval Function



```
void eval(char *cmdline)
{
    char *argv[MAXARGS]; /* Argument list execve() */
    char buf[MAXLINE];    /* Holds modified command line */
    int bg;               /* Should the job run in bg or fg? */
    pid_t pid;            /* Process id */

    strcpy(buf, cmdline);
    bg = parseline(buf, argv);
    if (argv[0] == NULL)
        return; /* Ignore empty lines */

    if (!builtin_command(argv)) {
        if ((pid = fork()) == 0) { /* Child runs user job */
            if (execve(argv[0], argv, environ) < 0) {
                printf("%s: Command not found.\n", argv[0]);
                exit(0);
            }
        }

        /* Parent waits for foreground job to terminate */
        if (!bg) {
            int status;
            if (waitpid(pid, &status, 0) < 0)
                unix_error("waitfg: waitpid error");
        }
        else
            printf("%d %s", pid, cmdline);
    }
    return;
}
```

哎呀。此代码有问题。  
Oops. There is a problem with this code.

# 简单Shell程序存在的问题

## Problem with Simple Shell Example



- **Shell设计成无限循环运行** Shell designed to run indefinitely
  - 不应该积累不需要的资源/Should not accumulate unneeded resources
    - 内存 Memory
    - 子进程 Child processes
    - 文件描述符 File descriptors
- **例子shell只能等待并回收前台作业** Our example shell correctly waits for and reaps foreground jobs
- **后台作业怎么办？** But what about background jobs?
  - 终止后变成僵尸 Will become zombies when they terminate
  - 由于shell不会终止，所以永远不会被回收 Will never be reaped because shell (typically) will not terminate
  - 会造成系统内存泄露并耗尽内核内存 Will create a memory leak that could run the kernel out of memory



# 可以利用ECF解决 ECF to the Rescue!

- **解决方案：异常控制流 Solution: Exceptional control flow**
  - 在后台进程处理完成后，内核打断正常处理流程并提醒我们 The kernel will interrupt regular processing to alert us when a background process completes
  - Unix系统中这种提醒的机制是信号 In Unix, the alert mechanism is called a *signal*





# 议题

- 外壳 Shells
- 信号 Signals
- 非局部跳转 Nonlocal jumps

# 信号 Signals



- 信号是一条小消息，用来通知一个进程某种类型的事件在系统中发生了 A **signal** is a small message that notifies a process that an event of some type has occurred in the system
  - 类似于异常和中断 Akin to exceptions and interrupts
  - 由内核发送给一个进程（有时是根据另一个进程的请求） Sent from the kernel (sometimes at the request of another process) to a process
  - 信号的类型是用1-30的小整型标识 Signal type is identified by small integer ID's (1-30)
  - 信号的唯一信息就是这个ID以及信号达到的事实 Only information in a signal is its ID and the fact that it arrived

<i>ID</i>	<i>Name</i>	<i>Default Action</i>	<i>Corresponding Event</i>
2	SIGINT	Terminate	用户输入ctrl-c User typed ctrl-c
9	SIGKILL	Terminate	杀死程序（不能覆盖或被忽略） Kill program (cannot override or ignore)
11	SIGSEGV	Terminate	段错误 Segmentation violation
14	SIGALRM	Terminate	时钟信号 Timer signal
17	SIGCHLD	Ignore	子进程停止或者终止 Child stopped or terminated



# 信号概念：发送一个信号

## Signal Concepts: Sending a Signal

- 内核通过更新目标进程上下文的某些状态来**发送**（传递）一个信号给**目标进程** Kernel **sends** (delivers) a signal to a **destination process** by updating some state in the context of the destination process
- 内核发送信号是由于以下原因之一 Kernel sends a signal for one of the following reasons:
  - 内核侦测到除零错误（SIGFPE）或者子进程终止（SIGCHLD）等系统事件 Kernel has detected a system event such as divide-by-zero (SIGFPE) or the termination of a child process (SIGCHLD)
  - 另外一个进程调用了**kill系统调用**显式请求内核发送一个信号给目标进程 Another process has invoked the **kill** system call to explicitly request the kernel to send a signal to the destination process

# 信号概念：发送一个信号

## Signal Concepts: Sending a Signal



用户级  
User level

进程B  
Process B

进程A  
Process A

进程C  
Process C

内核  
kernel

		挂起	Pending for A
		挂起	Pending for B
		挂起	Pending for C

		阻塞	Blocked for A
		阻塞	Blocked for B
		阻塞	Blocked for C

# 信号概念：发送一个信号

## Signal Concepts: Sending a Signal



用户级  
User level

进程B  
Process B

进程A  
Process A

进程C  
Process C

Sends to C

			Pending for A
			Pending for B
			Pending for C

			Blocked for A
			Blocked for B
			Blocked for C

内核  
kernel

# 信号概念：发送一个信号

## Signal Concepts: Sending a Signal



用户级  
User level

进程B  
Process B

进程A  
Process A

进程C  
Process C

内核  
kernel

			Pending for A
			Pending for B
		1	Pending for C

			Blocked for A
			Blocked for B
			Blocked for C

# 信号概念：发送一个信号Signal

## Concepts: Sending a Signal



用户级  
User level

进程B  
Process B

进程A  
Process A

进程C  
Process C

内核  
kernel

			Pending for A
			Pending for B
1			Pending for C

			Blocked for A
			Blocked for B
			Blocked for C

Received by C

# 信号概念：发送一个信号Signal

## Concepts: Sending a Signal



用户级  
User level

进程B  
Process B

进程A  
Process A

进程C  
Process C

内核  
kernel

			Pending for A
			Pending for B
		0	Pending for C

			Blocked for A
			Blocked for B
			Blocked for C

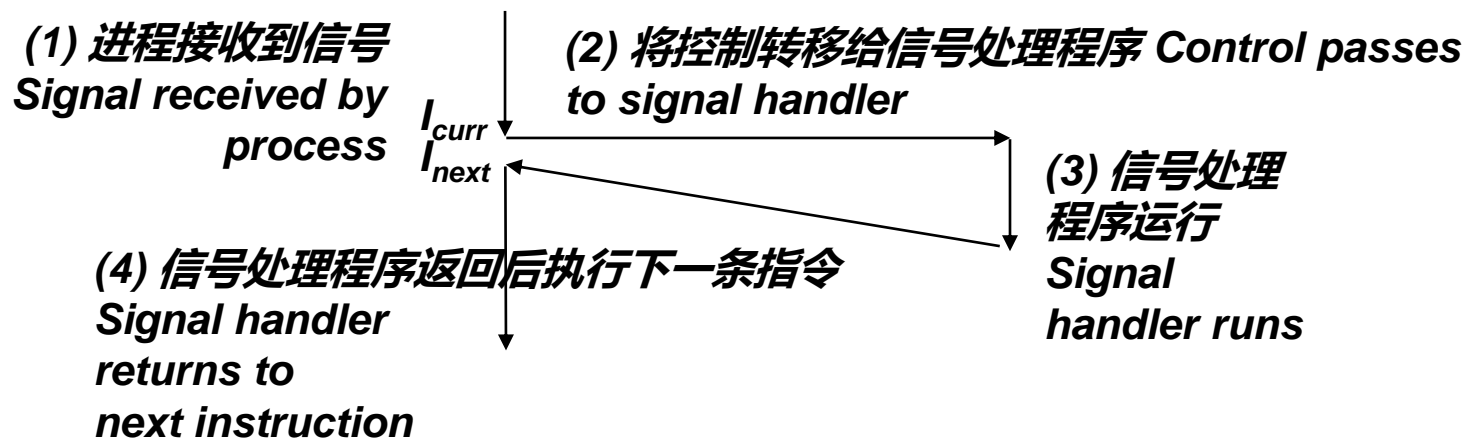




# 信号概念：接收一个信号

## Signal Concepts: Receiving a Signal

- 目标进程**接收**信号是由于系统内核强制其对某个信号的发送做出响应 A destination process **receives** a signal when it is forced by the kernel to react in some way to the delivery of the signal
- 可能的响应方式 Some possible ways to react:
  - **忽略**信号（什么也不做） **Ignore** the signal (do nothing)
  - **终止进程**（可以选择对信息转储） **Terminate** the process (with optional core dump)
  - **调用**用户级**信号处理函数**对信号进行处理 **Catch** the signal by executing a user-level function called **signal handler**
    - 类似于硬件异常处理函数对异步中断的响应 Akin to a hardware exception handler being called in response to an asynchronous interrupt:



# 信号概念：挂起或者阻塞的信号

## Signal Concepts: Pending and Blocked Signals



- 已经发送但是没有被接收的信号处于**挂起**状态 A signal is **pending** if sent but not yet received
  - 任何特定类型的信号最多有一个挂起的 There can be at most one pending signal of any particular type
  - 重要：信号不排队 Important: Signals are not queued
    - 如有某个进程有一个类型为k的信号挂起，则后续发给该进程的k类信号被直接抛弃 If a process has a pending signal of type k, then subsequent signals of type k that are sent to that process are discarded
- 一个进程会**阻塞**某种特定类型信号的接收 A process can **block** the receipt of certain signals
  - 阻塞的信号可以发送，但是在解除阻塞前不会被接收 Blocked signals can be delivered, but will not be received until the signal is unblocked
  - 有些信号不能被阻塞（SIGKILL, SIGSTOP）或者仅当其它进程发送（SIGSEGV、SIGILL等）时被阻塞 Some signals cannot be blocked (SIGKILL, SIGSTOP) or can only be blocked when sent by other processes (SIGSEGV, SIGILL, etc)
- 挂起的信号最多被接收一次 A pending signal is received at most once



# 信号概念：挂起/阻塞位

## Signal Concepts: Pending/Blocked Bits

- 内核在每个进程的上下文维护一个挂起和阻塞的比特向量 Kernel maintains `pending` and `blocked` bit vectors in the context of each process
  - 挂起：表示挂起的信号集合 `pending`: represents the set of pending signals
    - 当发送了一个k类型的信号时系统设置第k个比特位 Kernel sets bit k in `pending` when a signal of type k is delivered
    - 当类型k的信号被接收后系统会将第k个比特位清零 Kernel clears bit k in `pending` when a signal of type k is received
  - 阻塞：表示阻塞的信号集合 `blocked`: represents the set of blocked signals
    - 可以使用`sigprocmask`函数设置或者清除 Can be set and cleared by using the `sigprocmask` function
    - 也称为信号掩码 Also referred to as the *signal mask*.

# 信号概念：发送信号

## Signal Concepts: Sending a Signal



用户级  
User level

进程B  
Process B

进程A  
Process A

进程C  
Process C

Sends to C

内核  
kernel

			Pending for A
			Pending for B
	1		Pending for C

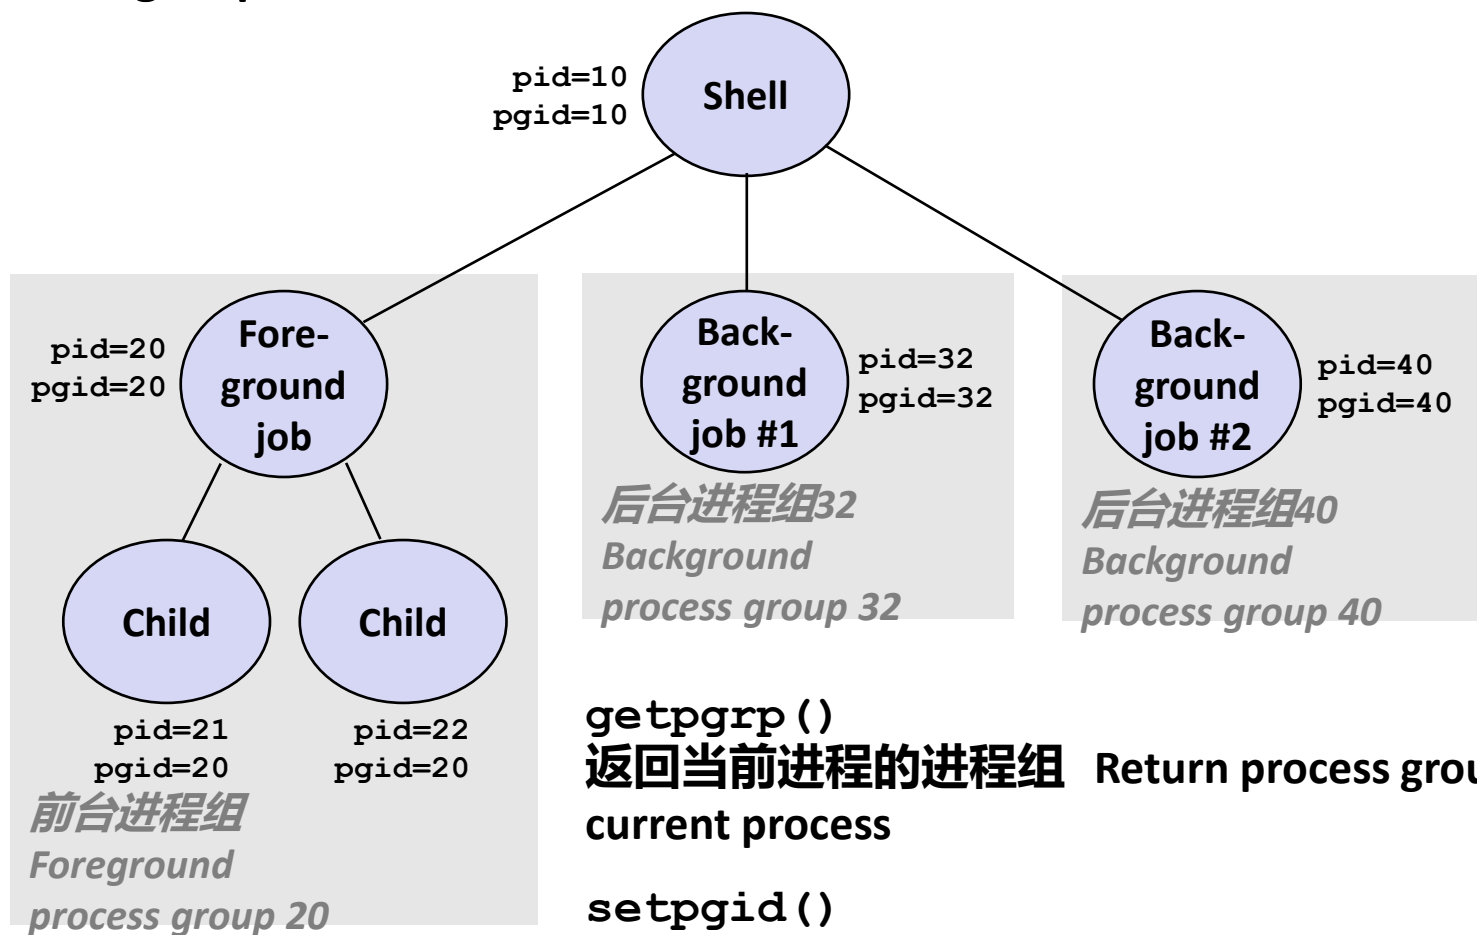
			Blocked for A
			Blocked for B
			Blocked for C



# 发送信号：进程组

## Sending Signals: Process Groups

- 每个进程只属于一个进程组 Every process belongs to exactly one process group



`getpgrp()`  
返回当前进程的进程组 Return process group of current process

`setpgid()`  
修改当前进程的进程组 (细节见教材) Change process group of a process (see text for details)



# 通过/bin/kill程序发送信号

## Sending Signals with /bin/kill Program

- /bin/kill程序可以发送任意信号给一个进程或者进程组 /bin/kill program sends arbitrary signal to a process or process group

- 例如 Examples

- /bin/kill -9 24818 发送SIGKILL给进程 24818 Send SIGKILL to process 24818
- /bin/kill -9 -24817 发送SIGKILL给进程组的每个进程 Send SIGKILL to every process in process group 24817

```
linux> ./forks 16
Child1: pid=24818 pgrp=24817
Child2: pid=24819 pgrp=24817
```

```
linux> ps
  PID TTY          TIME CMD
24788 pts/2        00:00:00 tcsh
24818 pts/2        00:00:02 forks
24819 pts/2        00:00:02 forks
24820 pts/2        00:00:00 ps
```

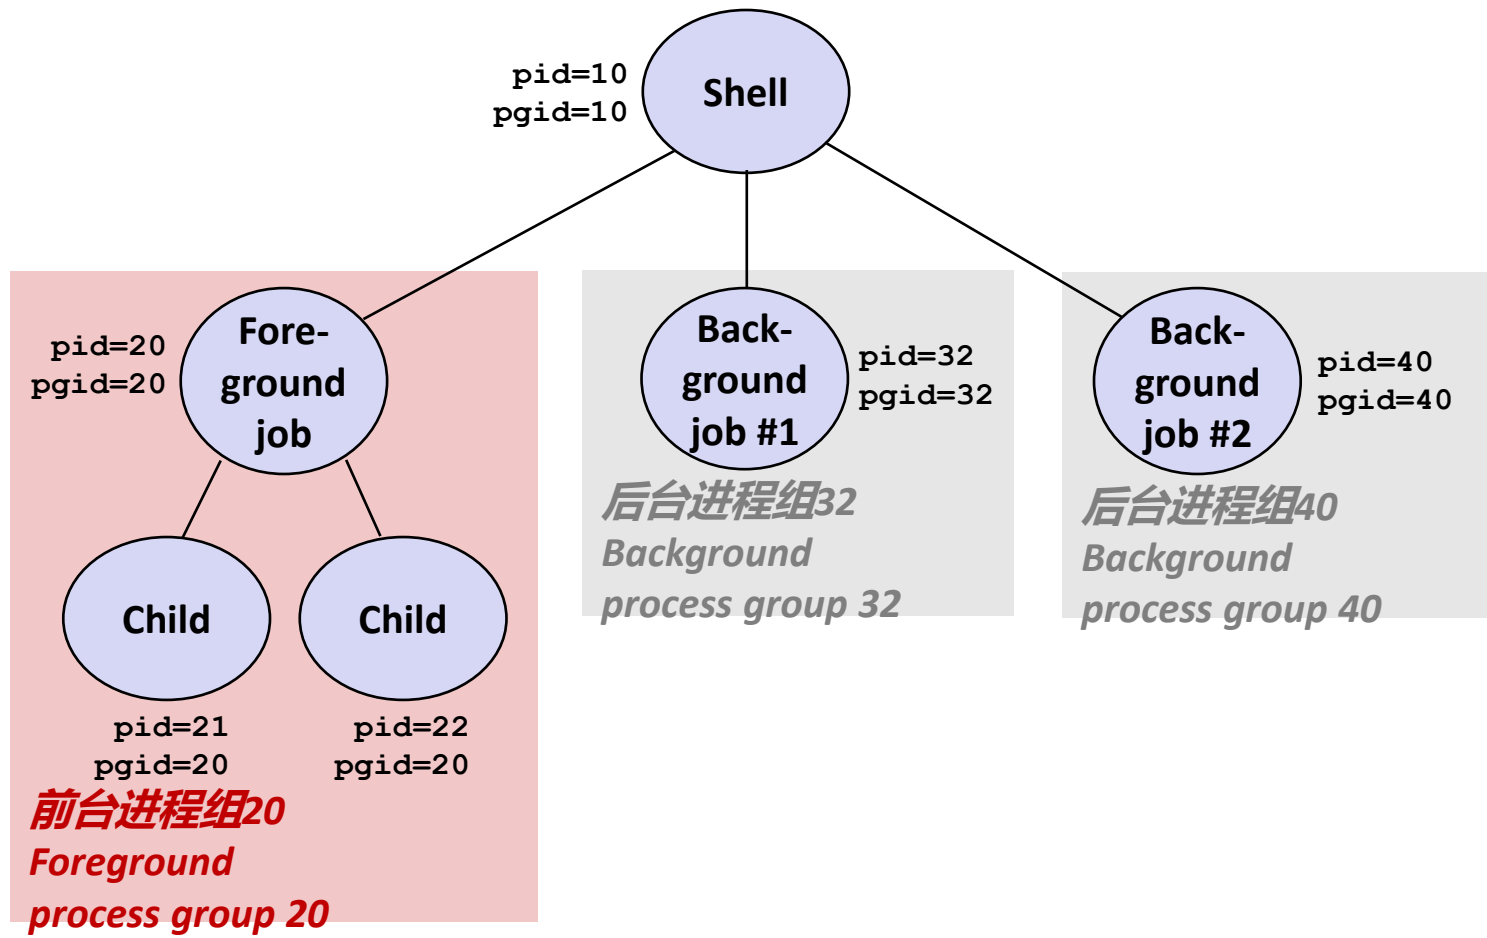
```
linux> /bin/kill -9 -24817
```

```
linux> ps
  PID TTY          TIME CMD
24788 pts/2        00:00:00 tcsh
24823 pts/2        00:00:00 ps
linux>
```

# 通过键盘发送信号 Sending Signals from the Keyboard



- 输入ctrl-c(ctrl-z)会导致系统内核发送一个SIGINT (SIGTSTP) 信号给前台进程组的每个作业 Typing ctrl-c (ctrl-z) causes the kernel to send a SIGINT (SIGTSTP) to every job in the foreground process group.
  - SIGINT – default action is to terminate each process 默认终止每个进程
  - SIGTSTP – default action is to stop (suspend) each process 默认停止（挂起）每个进程





# ctrl-c和ctrl-z示例

## Example of ctrl-c and ctrl-z

```
bluefish> ./forks 17
Child: pid=28108 pgrp=28107
Parent: pid=28107 pgrp=28107
<types ctrl-z>
Suspended
bluefish> ps w
  PID TTY          STAT       TIME COMMAND
 27699 pts/8        Ss          0:00   -tcsh
 28107 pts/8        T           0:01   ./forks 17
 28108 pts/8        T           0:01   ./forks 17
 28109 pts/8        R+          0:00   ps w
bluefish> fg
./forks 17
<types ctrl-c>
bluefish> ps w
  PID TTY          STAT       TIME COMMAND
 27699 pts/8        Ss          0:00   -tcsh
 28110 pts/8        R+          0:00   ps w
```

进程状态STAT标记 STAT  
(process state) Legend:

**First letter 第一个字母:**

S: sleeping 睡眠

T: stopped 停止

R: running 运行

**Second letter 第二个字母:**

s: session leader 会话首领

+: foreground proc group 前台  
进程组

参见“man ps”了解更多细节

See “man ps” for more  
details



# 通过kill函数发送信号

## Sending Signals with kill Function



```
void fork12()
{
    pid_t pid[N];
    int i;
    int child_status;

    for (i = 0; i < N; i++)
        if ((pid[i] = fork()) == 0) {
            /* Child: Infinite Loop */
            while(1)
                ;
        }

    for (i = 0; i < N; i++) {
        printf("Killing process %d\n", pid[i]);
        kill(pid[i], SIGINT);
    }

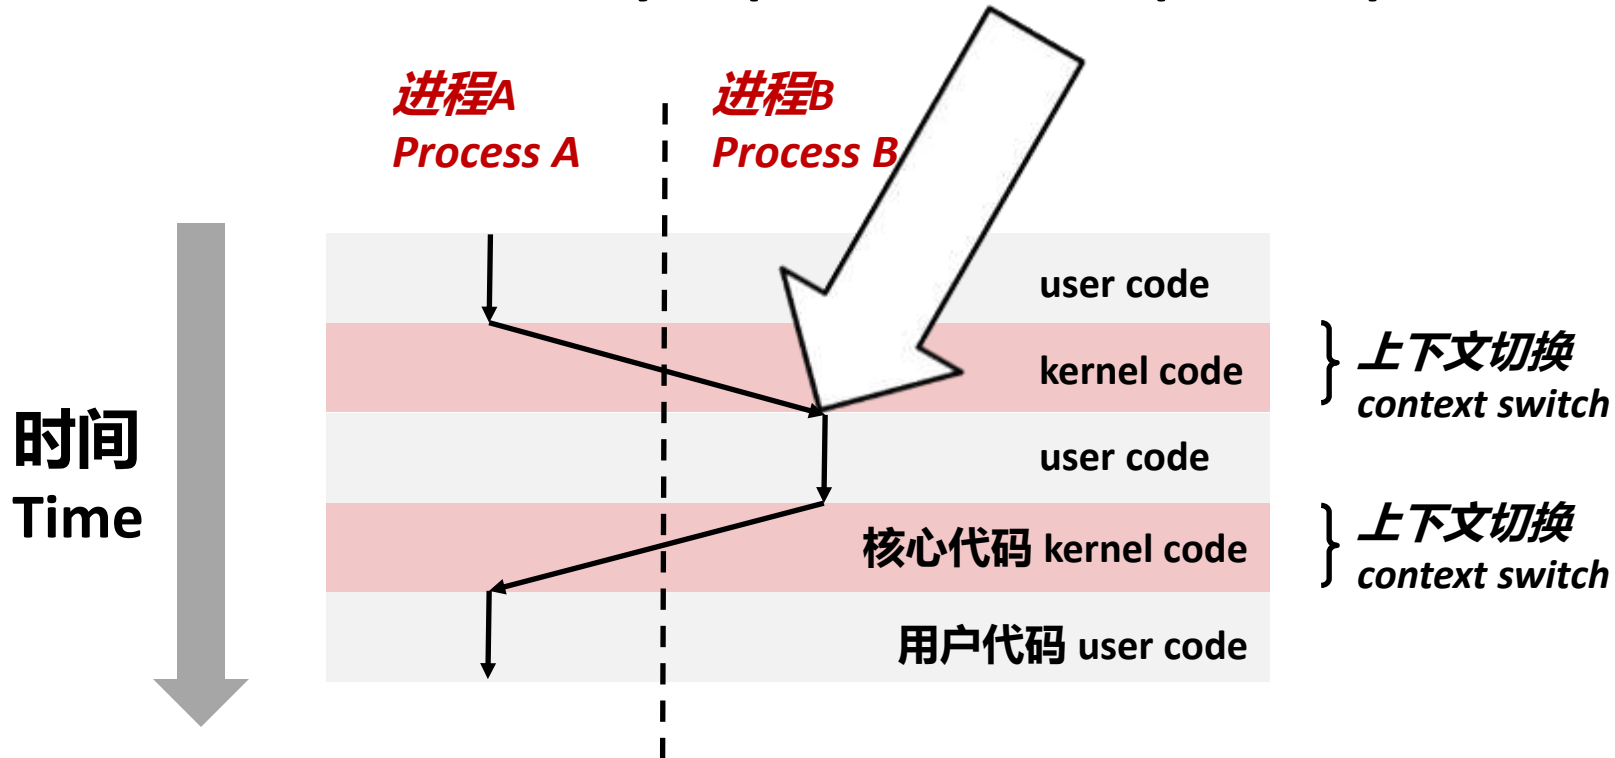
    for (i = 0; i < N; i++) {
        pid_t wpid = wait(&child_status);
        if (WIFEXITED(child_status))
            printf("Child %d terminated with exit status %d\n",
                wpid, WEXITSTATUS(child_status));
        else
            printf("Child %d terminated abnormally\n", wpid);
    }
}
```

*forks.c*



# 接收信号 Receiving Signals

- 假设内核正从异常处理函数返回，并准备把控制权传递给进程p  
Suppose kernel is returning from an exception handler and is ready to pass control to process  $p$



# 接收信号 Receiving Signals



- 假设内核正从异常处理函数返回，并准备把控制权传递给进程  $p$   
Suppose kernel is returning from an exception handler and is ready to pass control to process  $p$
- 内核计算 Kernel computes  $pnb = pending \ \& \ \sim blocked$ 
  - 进程  $p$  挂起但非阻塞信号的集合 The set of pending nonblocked signals for process  $p$
- 如果集合为空 If  $(pnb == 0)$ 
  - 将控制权交给进程  $p$  逻辑流的下一条指令 Pass control to next instruction in the logical flow for  $p$
- 否则 Else
  - 选择  $pnb$  中最低非0位  $k$  并强制进程  $p$  接收信号  $k$  Choose least nonzero bit  $k$  in  $pnb$  and force process  $p$  to **receive** signal  $k$
  - 信号的接收触发了  $p$  的某些动作 The receipt of the signal triggers some **action** by  $p$
  - 对  $pnb$  中每个非0位  $k$  重复上述过程 Repeat for all nonzero  $k$  in  $pnb$
  - 将控制权交给进程  $p$  逻辑流的下一条指令 Pass control to next instruction in logical flow for  $p$



# 默认动作 Default Actions

- 每种类型的信号有一个预定义的**默认动作**，可能是如下中的一个 Each signal type has a predefined **default action**, which is one of:
  - 终止进程 The process terminates
  - 停止进程，直到接收到SIGCONT时重启 The process stops until restarted by a SIGCONT signal
  - 进程忽略掉该信号 The process ignores the signal

# 安装信号处理程序 Installing Signal Handlers



- 函数signal修改接收信号signum对应的默认行为 The `signal` function modifies the default action associated with the receipt of signal `signum`:
  - `handler_t *signal(int signum, handler_t *handler)`
- 信号处理程序handler的不同值 Different values for `handler`:
  - `SIG_IGN`: ignore signals of type `signum` 忽略`signum`类型的信号
  - `SIG_DFL`: revert to the default action on receipt of signals of type `signum` 接收到`signum`类型的信号时按照默认动作处理
  - 否则`handler`是用户级信号处理程序的地址 Otherwise, `handler` is the address of a user-level **signal handler**
    - 当进程接收到类型为`signum`的信号时调用 Called when process receives signal of type `signum`
    - 称为安装信号处理程序 Referred to as **“installing”** the handler
    - 执行信号处理程序称为捕获或处理该信号 Executing handler is called **“catching”** or **“handling”** the signal
    - 当信号处理程序执行返回语句时，控制权交给进程接收到信号时被打断控制流中指令 When the handler executes its return statement, control passes back to instruction in the control flow of the process that was interrupted by receipt of the signal

# 信号处理例子 Signal Handling Example



```
void sigint_handler(int sig) /* SIGINT handler */
{
    printf("So you think you can stop the bomb with ctrl-c, do you?\n");
    sleep(2);
    printf("Well...");
    fflush(stdout);
    sleep(1);
    printf("OK. :-)\n");
    exit(0);
}

int main()
{
    /* Install the SIGINT handler */
    if (signal(SIGINT, sigint_handler) == SIG_ERR)
        unix_error("signal error");

    /* Wait for the receipt of a signal */
    pause();

    return 0;
}
```

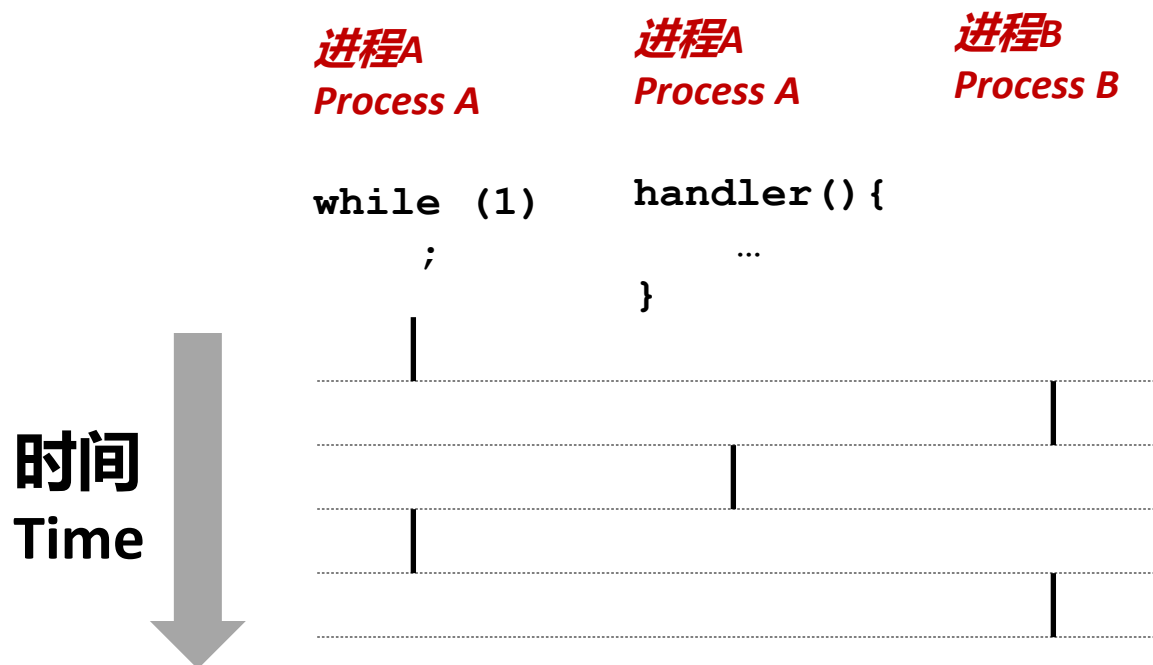
sigint.c

# 信号处理程序作为并发控制流

## Signals Handlers as Concurrent Flows



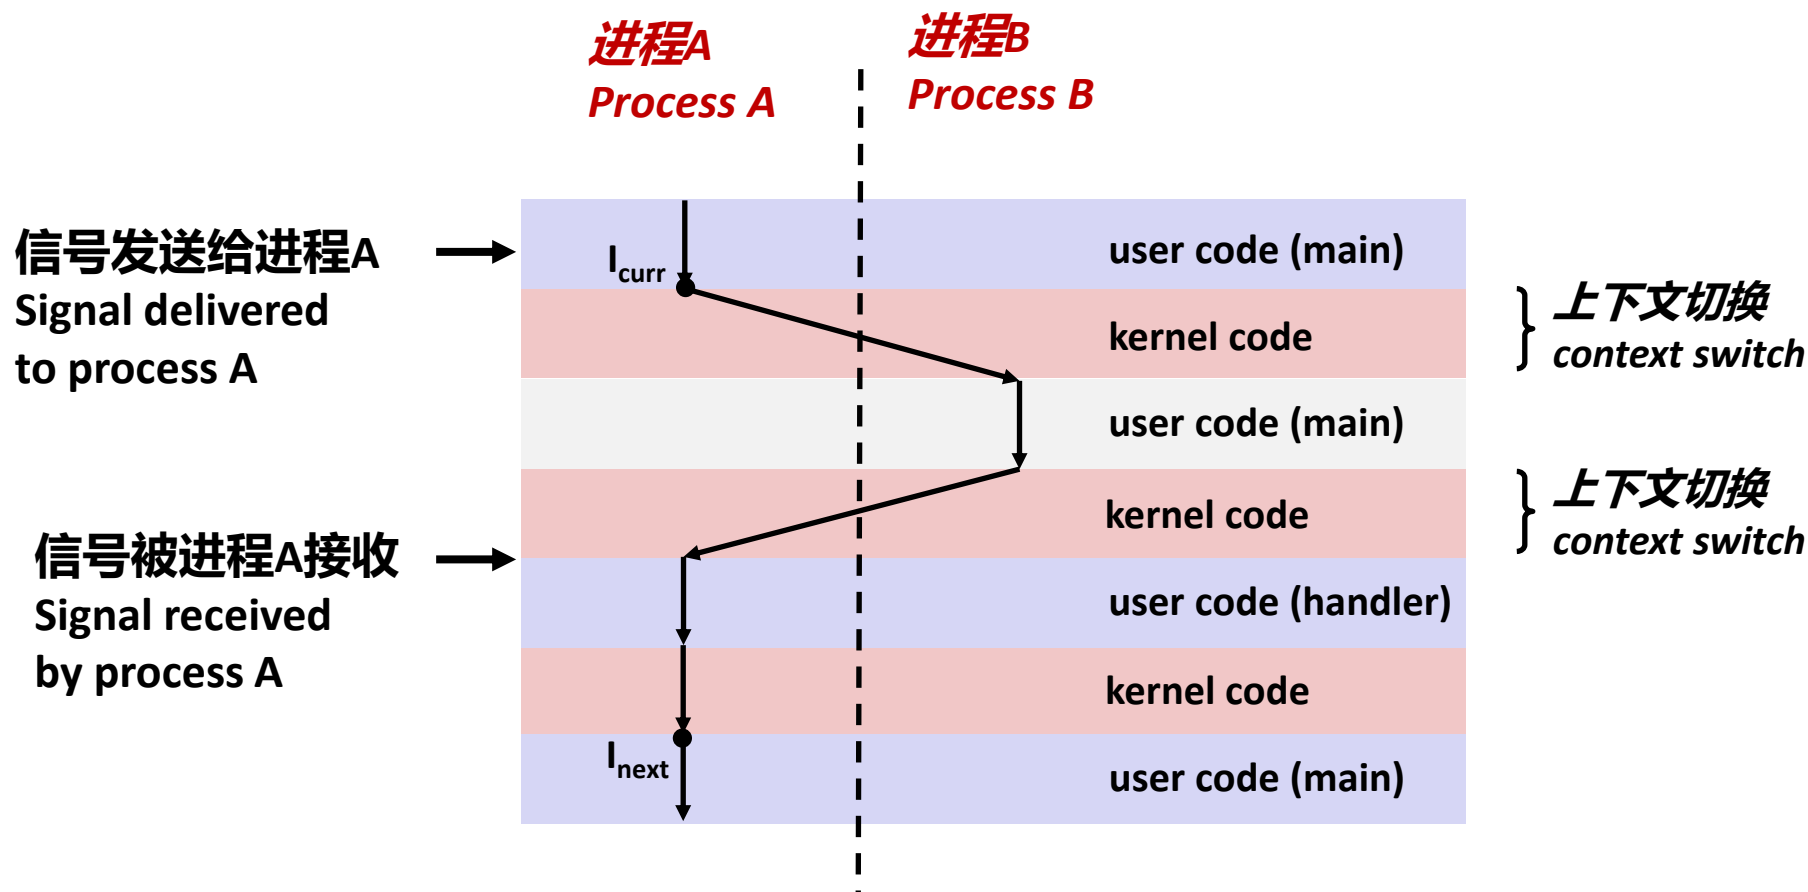
- 每个信号处理程序都是一个独立的逻辑控制流（非进程），与主程序并发执行 A signal handler is a separate logical flow (not process) that runs concurrently with the main program





# 信号处理程序作为并发控制流的另一个视图

## Another View of Signal Handlers as Concurrent Flows

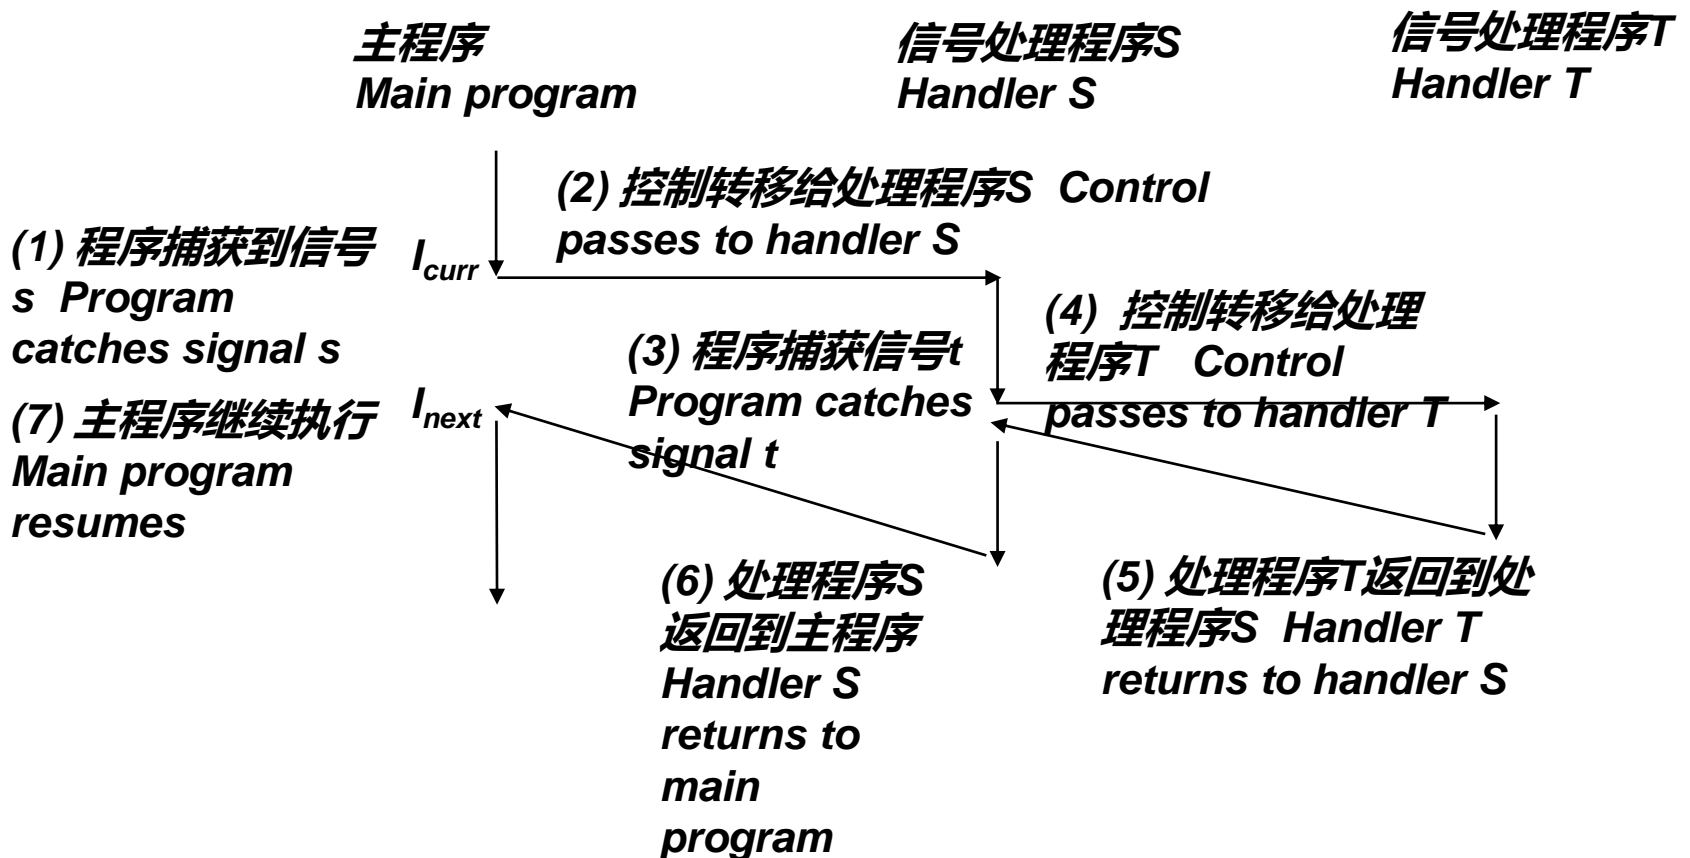






# 嵌套信号处理 Nested Signal Handlers

- 信号处理程序可能被另一个信号处理程序打断 Handlers can be interrupted by other handlers





# 阻塞和解除信号阻塞

## Blocking and Unblocking Signals

- **隐式阻塞机制 Implicit blocking mechanism**
  - 内核会阻塞当前正在被处理的任何挂起信号类型 Kernel blocks any pending signals of type currently being handled.
  - 例如SIGINT信号处理程序不能被另一个SIGINT打断 E.g., A SIGINT handler can't be interrupted by another SIGINT
- **显式阻塞和解除阻塞机制 Explicit blocking and unblocking mechanism**
  - `sigprocmask`函数 `sigprocmask` function
- **支持函数 Supporting functions**
  - `sigemptyset` – Create empty set 创建一个空的集合
  - `sigfillset` – Add every signal number to set 对集合设置每个信号编号
  - `sigaddset` – Add signal number to set 对集合设置某个信号编号
  - `sigdelset` – Delete signal number from set 将信号编号从集合删除

# 临时阻塞信号

## Temporarily Blocking Signals



```
sigset_t mask, prev_mask;
```

```
Sigemptyset(&mask);
```

```
Sigaddset(&mask, SIGINT);
```

```
/* Block SIGINT and save previous blocked set */
```

```
Sigprocmask(SIG_BLOCK, &mask, &prev_mask);
```

```
⋮  
/* Code region that will not be interrupted by SIGINT */
```

```
/* Restore previous blocked set, unblocking SIGINT */
```

```
Sigprocmask(SIG_SETMASK, &prev_mask, NULL);
```

# 安全的信号处理

## Safe Signal Handling



- **信号处理程序比较复杂，是因为他们是和主程序并发运行的，并且共享同样的全局数据结构** Handlers are tricky because they are concurrent with main program and share the same global data structures.
  - **共享数据结构更容易被破坏** Shared data structures can become corrupted.
- **我们在这学期后面讨论并发的问题** We'll explore concurrency issues later in the term.
- **现在只给一些有助避免麻烦的提示** For now here are some guidelines to help you avoid trouble.

# 编写安全处理程序的提示

## Guidelines for Writing Safe Handlers



- **G0: 信号处理程序越简单越好** Keep your handlers as simple as possible
  - 例如, 设置全局标记后返回 e.g., Set a global flag and return
- **G1: 在信号处理程序中只调用异步信号安全的函数** Call only async-signal-safe functions in your handlers
  - `printf`, `sprintf`, `malloc`, and `exit` are not safe! 这些都不安全
- **G2: 进入和退出时保存和恢复 `errno`** Save and restore `errno` on entry and exit
  - 以便其它的信号处理程序不会覆盖你的 `errno` 值 So that other handlers don't overwrite your value of `errno`
- **G3: 临时阻塞所有的信号后再访问共享数据结构** Protect accesses to shared data structures by temporarily blocking all signals.
  - 避免可能的破坏 To prevent possible corruption
- **G4: 将全局变量声明为 `volatile`** Declare global variables as `volatile`
  - 避免编译器将其存储在寄存器中 To prevent compiler from storing them in a register
- **G5: 将全局标记声明为 `volatile sig_atomic_t`** Declare global flags as `volatile sig_atomic_t`
  - `flag` 只读或只写的变量 (例如 `flag=1`, 不是 `flag++`) `flag`: variable that is only read or written (e.g. `flag = 1`, not `flag++`)
  - 按照这种方式声明的 `flag` 变量不需要像其他全局变量那样保护 Flag declared this way does not need to be protected like other globals

# 异步信号安全 Async-Signal-Safety



- 如果一个函数是可重入的（例如所有变量存储在栈帧，CS:APP3e 12.7.2）或者不可以被信号打断的则将其称为**异步信号安全***async-signal-safe* Function is *async-signal-safe* if either reentrant (e.g., all variables stored on stack frame, CS:APP3e 12.7.2) or non-interruptible by signals.
- Posix中有117个函数是异步信号安全*async-signal-safe* Posix guarantees 117 functions to be *async-signal-safe*
  - 来源：man命令 Source: “man 7 signal”
  - 在其中的常见函数包括： Popular functions on the list:
    - `_exit`, `write`, `wait`, `waitpid`, `sleep`, `kill`
  - 常见的函数并不在其中 Popular functions that are **not** on the list:
    - `printf`, `sprintf`, `malloc`, `exit`
    - 不幸的事实： `write`是唯一异步信号安全*async-signal-safe*输出函数  
Unfortunate fact: `write` is the only *async-signal-safe* output function

# 安全格式化输出：选项#1



## Safe Formatted Output: Option #1

- 在信号处理程序中使用csapp.c的可重入的SIO（安全I/O库）

Use the reentrant SIO (Safe I/O library) from `csapp.c` in your handlers

- `ssize_t sio_puts(char s[]) /* Put string */`
- `ssize_t sio_putl(long v) /* Put long */`
- `void sio_error(char s[]) /* Put msg & exit */`

```
void sigint_handler(int sig) /* Safe SIGINT handler */
{
    sio_puts("So you think you can stop the bomb"
             " with ctrl-c, do you?\n");
    sleep(2);
    sio_puts("Well...");
    sleep(1);
    sio_puts("OK. :-)\n");
    _exit(0);
}
```

sigintsafe.c

# 安全格式化输出：选项#2



## Safe Formatted Output: Option #2

- 使用新的且改进的可重入 `sio_printf`! Use the new & improved reentrant `sio_printf`!
  - 处理 `printf` 受限类的格式串 Handles restricted class of `printf` format strings
    - 识别: Recognizes: `%c %s %d %u %x %%`
    - 大小指定符: Size designators `'l'` and `'z'`

```
void sigint_handler(int sig) /* Safe SIGINT handler */
{
    sio_printf("So you think you can stop the bomb"
               " (process %d) with ctrl-%c, do you?\n",
               (int) getpid(), 'c');

    sleep(2);
    sio_puts("Well...");
    sleep(1);
    sio_puts("OK. :-)\n");
    _exit(0);
}
```

sigintsafe.c



# 正确的信号处理

## Correct Signal Handling

```
volatile int ccount = 0;
void child_handler(int sig) {
    int olderrno = errno;
    pid_t pid;
    if ((pid = wait(NULL)) < 0)
        Sio_error("wait error");
    ccount--;
    sio_puts("Handler reaped child ");
    sio_putl((long)pid);
    sio_puts(" \n");
    sleep(1);
    errno = olderrno;
}

void fork14() {
    pid_t pid[N];
    int i;
    ccount = N;
    signal(SIGCHLD, child_handler);

    for (i = 0; i < N; i++) {
        if ((pid[i] = fork()) == 0) {
            sleep(1);
            exit(0); /* Child exits */
        }
    }
    while (ccount > 0) /* Parent spins */
        ;
}
```

**这段代码不正确!**  
**This code is incorrect!**

**N == 5**

### ■ 挂起的信号是不排队的 Pending signals are not queued

- 对每个信号类型，只用一个比特位来标识是否有信号被挂起 For each signal type, one bit indicates whether or not signal is pending...
- 因此每种最多有一个挂起的信号 ...thus at most one pending signal of any particular type.

### ■ 不可以使用信号对事件计数，例如子进程终止等 You can't use signals

```
whaleshark> ./forks 14
Handler reaped child 23240
Handler reaped child 23241
...(hangs)
```

as

# 正确信号处理 Correct Signal Handling



- 必须等待所有终止的子进程 Must wait for all terminated child processes
  - 将wait放入到循环中以回收所有终止的子进程 Put `wait` in a loop to reap all terminated children

```
void child_handler2(int sig)
{
    int olderrno = errno;
    pid_t pid;
    while ((pid = wait(NULL)) > 0) {
        ccount--;
        sio_puts("Handler reaped child ");
        sio_putl((long)pid);
        sio_puts(" \n");
    }
    if (errno != ECHILD)
        sio_error("wait error");
    errno = olderrno;
}
```

```
whaleshark> ./forks 15
Handler reaped child 23246
Handler reaped child 23247
Handler reaped child 23248
Handler reaped child 23249
Handler reaped child 23250
whaleshark>
```

# 可移植的信号处理

## Portable Signal Handling



- **不同的Unix版本有不同的信号处理语义**    Ugh! Different versions of Unix can have different signal handling semantics
  - 一些早期的系统在捕获到信号后会恢复默认动作    Some older systems restore action to default after catching signal
  - 有些被中断的系统调用会返回 `errno == EINTR`    Some interrupted system calls can return with `errno == EINTR`
  - 有的系统并不阻塞正在被处理的信号类型    Some systems don't block signals of the type being handled
- **解决方案: `sigaction`**    Solution: `sigaction`

```
handler_t *Signal(int signum, handler_t *handler)
{
    struct sigaction action, old_action;

    action.sa_handler = handler;
    sigemptyset(&action.sa_mask); /* Block sigs of type being handled */
    action.sa_flags = SA_RESTART; /* Restart syscalls if possible */

    if (sigaction(signum, &action, &old_action) < 0)
        unix_error("Signal error");
    return (old_action.sa_handler);
}
```

csapp.c

# 同步控制流避免竞争

## Synchronizing Flows to Avoid Races



- 简单shell的SIGCHLD处理程序 SIGCHLD handler for a simple shell
  - 当运行临界代码时阻塞所有信号 Blocks all signals while running critical code

```
void handler(int sig)
{
    int olderrno = errno;
    sigset_t mask_all, prev_all;
    pid_t pid;

    sigfillset(&mask_all);
    while ((pid = waitpid(-1, NULL, 0)) > 0) { /* Reap child */
        sigprocmask(SIG_BLOCK, &mask_all, &prev_all);
        deletejob(pid); /* Delete the child from the job list */
        sigprocmask(SIG_SETMASK, &prev_all, NULL);
    }
    if (pid != 0 && errno != ECHILD)
        sio_error("waitpid error");
    errno = olderrno;
}
```

procmask1.c

# 同步控制流避免竞争



## Synchronizing Flows to Avoid Races

- 简单的shell程序有个不易发现的同步问题，因为其假设父进程先于子进程  
Simple shell with a subtle synchronization error because it assumes parent runs before child

```
int main(int argc, char **argv)
{
    int pid;
    sigset_t mask_all, prev_all;
    int n = N; /* N = 5 */
    sigfillset(&mask_all);
    signal(SIGCHLD, handler);
    initjobs(); /* Initialize the job list */

    while (n--) {
        if ((pid = fork()) == 0) { /* Child */
            execve("/bin/date", argv, NULL);
        }
        sigprocmask(SIG_BLOCK, &mask_all, &prev_all); /* Parent */
        addjob(pid); /* Add the child to the job list */
        sigprocmask(SIG_SETMASK, &prev_all, NULL);
    }
    exit(0);
}
```

procmask1.c

# 没有竞争问题的修正shell程序

## Corrected Shell Program Without Race



```
int main(int argc, char **argv)
{
    int pid;
    sigset_t mask_all, mask_one, prev_one;
    int n = N; /* N = 5 */
    sigfillset(&mask_all);
    sigemptyset(&mask_one);
    sigaddset(&mask_one, SIGCHLD);
    signal(SIGCHLD, handler);
    initjobs(); /* Initialize the job list */

    while (n--) {
        sigprocmask(SIG_BLOCK, &mask_one, &prev_one); /* Block SIGCHLD */
        if ((pid = fork()) == 0) { /* Child process */
            sigprocmask(SIG_SETMASK, &prev_one, NULL); /* Unblock SIGCHLD */
            execve("/bin/date", argv, NULL);
        }
        sigprocmask(SIG_BLOCK, &mask_all, NULL); /* Parent process */
        addjob(pid); /* Add the child to the job list */
        sigprocmask(SIG_SETMASK, &prev_one, NULL); /* Unblock SIGCHLD */
    }
    exit(0);
}
```

# 显式等待信号

## Explicitly Waiting for Signals



- 信号处理程序显式等待SIGCHLD信号的到来 Handlers for program explicitly waiting for SIGCHLD to arrive

```
volatile sig_atomic_t pid;

void sigchld_handler(int s)
{
    int olderrno = errno;
    pid = waitpid(-1, NULL, 0); /* Main is waiting for nonzero pid */
    errno = olderrno;
}

void sigint_handler(int s)
{
}
```

waitforsignal.c

# 显式等待信号 Explicitly Waiting for Signals



类似于shell等待一个前台的作业终止  
Similar to a shell waiting  
for a foreground job to terminate.

```
int main(int argc, char **argv) {
    sigset_t mask, prev;
    int n = N; /* N = 10 */
    signal(SIGCHLD, sigchld_handler);
    signal(SIGINT, sigint_handler);
    sigemptyset(&mask);
    sigaddset(&mask, SIGCHLD);

    while (n--) {
        sigprocmask(SIG_BLOCK, &mask, &prev); /* Block SIGCHLD */
        if (fork() == 0) /* Child */
            exit(0);
        /* Parent */
        pid = 0;
        sigprocmask(SIG_SETMASK, &prev, NULL); /* Unblock SIGCHLD */

        /* Wait for SIGCHLD to be received (wasteful!) */
        while (!pid)
            ;

        /* Do some work after receiving SIGCHLD */
        printf(".");
    }
    printf("\n");
    exit(0);
}
```



# 显式等待信号 Explicitly Waiting for Signals



```
while (!pid)
    ;
```

- **程序是对的，但是太浪费资源** Program is correct, but very wasteful

- 程序忙于等待循环 Program in busy-wait loop

```
while (!pid) /* Race! */
    pause();
```

- **可能存在竞争** Possible race condition

- 在检查pid和开始暂停之间，可能接收信号 Between checking pid and starting pause, might receive signal

```
while (!pid) /* Too slow! */
    sleep(1);
```

- **安全，但是很慢** Safe, but slow

- 会占用1秒钟才能响应 Will take up to one second to respond

- **Solution: sigsuspend**

# 使用sigsuspend等待信号

## Waiting for Signals with sigsuspend



- `int sigsuspend(const sigset_t *mask)`
- 等价于原子版本（无中断可能）的： Equivalent to atomic (uninterruptable) version of:

```
sigprocmask(SIG_SETMASK, &mask, &prev);  
pause();  
sigprocmask(SIG_SETMASK, &prev, NULL);
```

# 使用sigsuspend等待信号

## Waiting for Signals with sigsuspend



```
int main(int argc, char **argv) {
    sigset_t mask, prev;
    int n = N; /* N = 10 */
    signal(SIGCHLD, sigchld_handler);
    signal(SIGINT, sigint_handler);
    sigemptyset(&mask);
    sigaddset(&mask, SIGCHLD);
    while (n--) {
        sigprocmask(SIG_BLOCK, &mask, &prev); /* Block SIGCHLD */
        if (fork() == 0) /* Child */
            exit(0);

        /* Wait for SIGCHLD to be received */
        pid = 0;
        while (!pid)
            sigsuspend(&prev);
        /* Optionally unblock SIGCHLD */
        sigprocmask(SIG_SETMASK, &prev, NULL);
        /* Do some work after receiving SIGCHLD */
        printf(".");
    }
    printf("\n");
    exit(0);
}
```



# 议题

- 外壳 Shells
- 信号 Signals
- **非局部跳转 Nonlocal jumps**
  - 参见教材和附加的幻灯片 Consult your textbook and additional slides



# 总结 Summary

- **信号提供进程级异常处理** Signals provide process-level exception handling
  - 可以从用户程序产生 Can generate from user programs
  - 可以声明信号处理程序定义处理效果 Can define effect by declaring signal handler
  - 编写信号处理函数的时候要特别小心 Be very careful when writing signal handlers
- **非局部跳转给出了进程内部的异常控制流** Nonlocal jumps provide exceptional control flow within process
  - 遵守栈相关的原则 Within constraints of stack discipline



# 非局部跳转

## Nonlocal Jumps: `setjmp/longjmp`

- 将控制转移到任意位置的强大（但比较危险）用户级机制 Powerful (but dangerous) user-level mechanism for transferring control to an arbitrary location
  - 受控的打破call/return规则的方式 Controlled way to break the procedure call / return discipline
  - 通常用于错误恢复和信号处理 Useful for error recovery and signal handling
- `int setjmp(jmp_buf j)`
  - 必须在longjmp之前调用 Must be called before longjmp
  - 给出后续longjmp对应的返回位置 Identifies a return site for a subsequent longjmp
  - 一次调用，返回一次或者多次 Called **once**, returns **one or more** times
- 实现 Implementation:
  - 通过将当前寄存器上下文、栈指针和PC值存储在jmp\_buf中记住当前位置 Remember where you are by storing the current **register context**, **stack pointer**, and **PC value** in `jmp_buf`
  - 返回0 Return 0

# setjmp/longjmp (续 cont)



## ■ void longjmp(jmp\_buf j, int i)

### ■ 含义 Meaning:

- 从setjmp返回, 再次被跳转缓冲区j记住 return from the **setjmp** remembered by jump buffer **j** again ...
- 这次返回i而不是0 ... this time returning **i** instead of 0

### ■ setjmp之后调用 Called after **setjmp**

- 一次调用但是从不返回 Called **once**, but **never** returns

## ■ longjmp实现 longjmp Implementation:

- 从跳转缓冲区j中恢复寄存器上下文 (栈指针、基指针、PC值)  
Restore register context (stack pointer, base pointer, PC value) from jump buffer **j**
- 将返回值寄存器%eax设置为i Set **%eax** (the return value) to **i**
- 跳转到跳转缓冲j中PC指定的位置 Jump to the location indicated by the PC stored in jump buf **j**



# setjmp/longjmp Example 示例

- 目标：从深度嵌套的函数直接返回最开始的调用者
- Goal: return directly to original caller from a deeply-nested function

```
/* Deeply nested function foo */  
void foo(void)  
{  
    if (error1)  
        longjmp(buf, 1);  
    bar();  
}  
  
void bar(void)  
{  
    if (error2)  
        longjmp(buf, 2);  
}
```





```
jmp_buf buf;  
  
int error1 = 0;  
int error2 = 1;  
  
void foo(void), bar(void);
```

```
int main()  
{  
    switch(setjmp(buf)) {  
        case 0:  
            foo();  
            break;  
        case 1:  
            printf("Detected an error1 condition in foo\n");  
            break;  
        case 2:  
            printf("Detected an error2 condition in foo\n");  
            break;  
        default:  
            printf("Unknown error condition in foo\n");  
    }  
    exit(0);  
}
```

## setjmp/longjmp 示例/Example (续/cont)

# 非局部跳转的限制

## Limitations of Nonlocal Jumps



### ■ 基于栈原理工作 Works within stack discipline

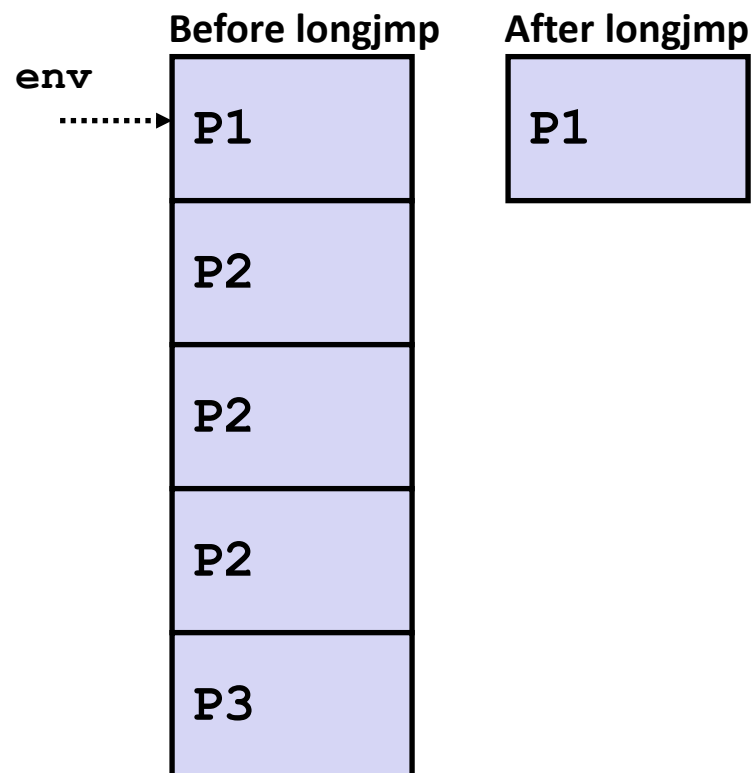
- 只能跳转到已经调用但是还没有完成的函数 Can only long jump to environment of function that has been called but not yet completed

```
jmp_buf env;

P1()
{
    if (setjmp(env)) {
        /* Long Jump to here */
    } else {
        P2();
    }
}

P2()
{
    . . . P2(); . . . P3();
}

P3()
{
    longjmp(env, 1);
}
```








# 非局部跳转的限制 (续)

## Limitations of Long Jumps (cont.)

### ■ 基于栈原理工作 Works within stack discipline

- 只能跳转到已经调用但是还没有完成的函数/Can only long jump to environment of function that has been called but not yet completed

```
jmp_buf env;  
  
P1()  
{  
    P2(); P3();  
}  
  
P2()  
{  
    if (setjmp(env)) {  
        /* Long Jump to here */  
    }  
}  
  
P3()  
{  
    longjmp(env, 1);  
}
```



# 整合在一起：程序在按下ctrl-c或d时重启



## Putting It All Together: A Program That Restarts Itself When ctrl-c'd

```
#include "csapp.h"

sigjmp_buf buf;

void handler(int sig)
{
    siglongjmp(buf, 1);
}

int main()
{
    if (!sigsetjmp(buf, 1)) {
        Signal(SIGINT, handler);
        Sio_puts("starting\n");
    }
    else
        Sio_puts("restarting\n");

    while(1) {
        Sleep(1);
        Sio_puts("processing...\n");
    }
    exit(0); /* Control never reaches here */
}
```

```
greatwhite> ./restart
starting
processing...
processing...
processing...
restarting
processing... ← Ctrl-c
processing...
restarting
processing... ← Ctrl-c
processing...
processing...
```

restart.c